



# **Computer Graphics - Filtering**

Junjie Cao @ DLUT

Spring 2016

<http://jjcao.github.io/ComputerGraphics/>

# Denoising





# Blur







# Sharpen





# Edge detection





# A “smarter” blur (doesn't blur over edges)



# Bilateral filter





# Convolution

$$(f * g)(x) = \int_{-\infty}^{\infty} \underbrace{f(y)}_{\text{filter}} \underbrace{g(x-y)}_{\text{input signal}} dy$$


output signal

- It may be helpful to consider the effect of convolution with the simple unit-area “box” function:



$$f(x) = \begin{cases} 1 & |x| \leq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

$$(f * g)(x) = \int_{-0.5}^{0.5} g(x-y) dy$$

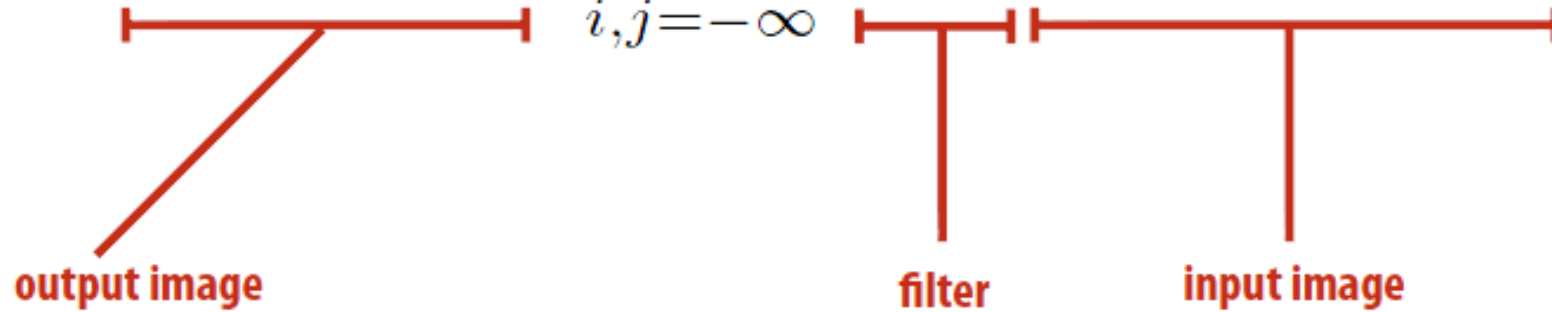
$f * g$  is a “smoothed” version of  $g$







# Discrete 2D convolution

$$(f * g)(x, y) = \sum_{i, j=-\infty}^{\infty} f(i, j) I(x - i, y - j)$$


A diagram illustrating the components of the convolution equation. A horizontal line represents the summation over  $i, j$ . A red bracket under the summation index  $i, j$  is labeled "filter". A red bracket under the input image  $I(x - i, y - j)$  is labeled "input image". A red line points from the output  $(f * g)(x, y)$  to the label "output image".


Consider  $f(i, j)$  that is nonzero only when:  $-1 \leq i, j \leq 1$

Then:

$$(f * g)(x, y) = \sum_{i, j=-1}^1 f(i, j) I(x - i, y - j)$$


And we can represent  $f(i, j)$  as a 3x3 matrix of values where:

$$f(i, j) = \mathbf{F}_{i, j} \quad \text{(often called: "filter weights", "kernel")}$$






# Linear filtering (warm-up slide)




original





?

# Linear filtering (warm-up slide)




original




Filtered  
(no change)



# Linear filtering




original


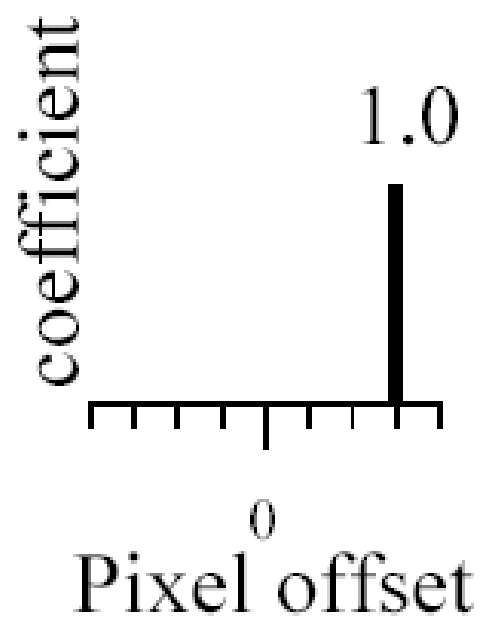


?

shift




original




shifted



# Linear filtering




original


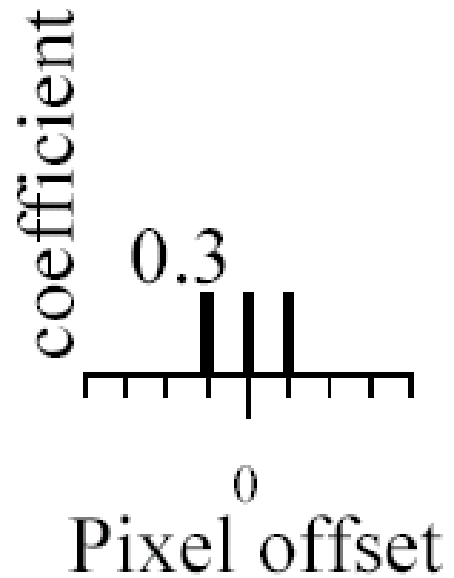


?

# Blurring




original



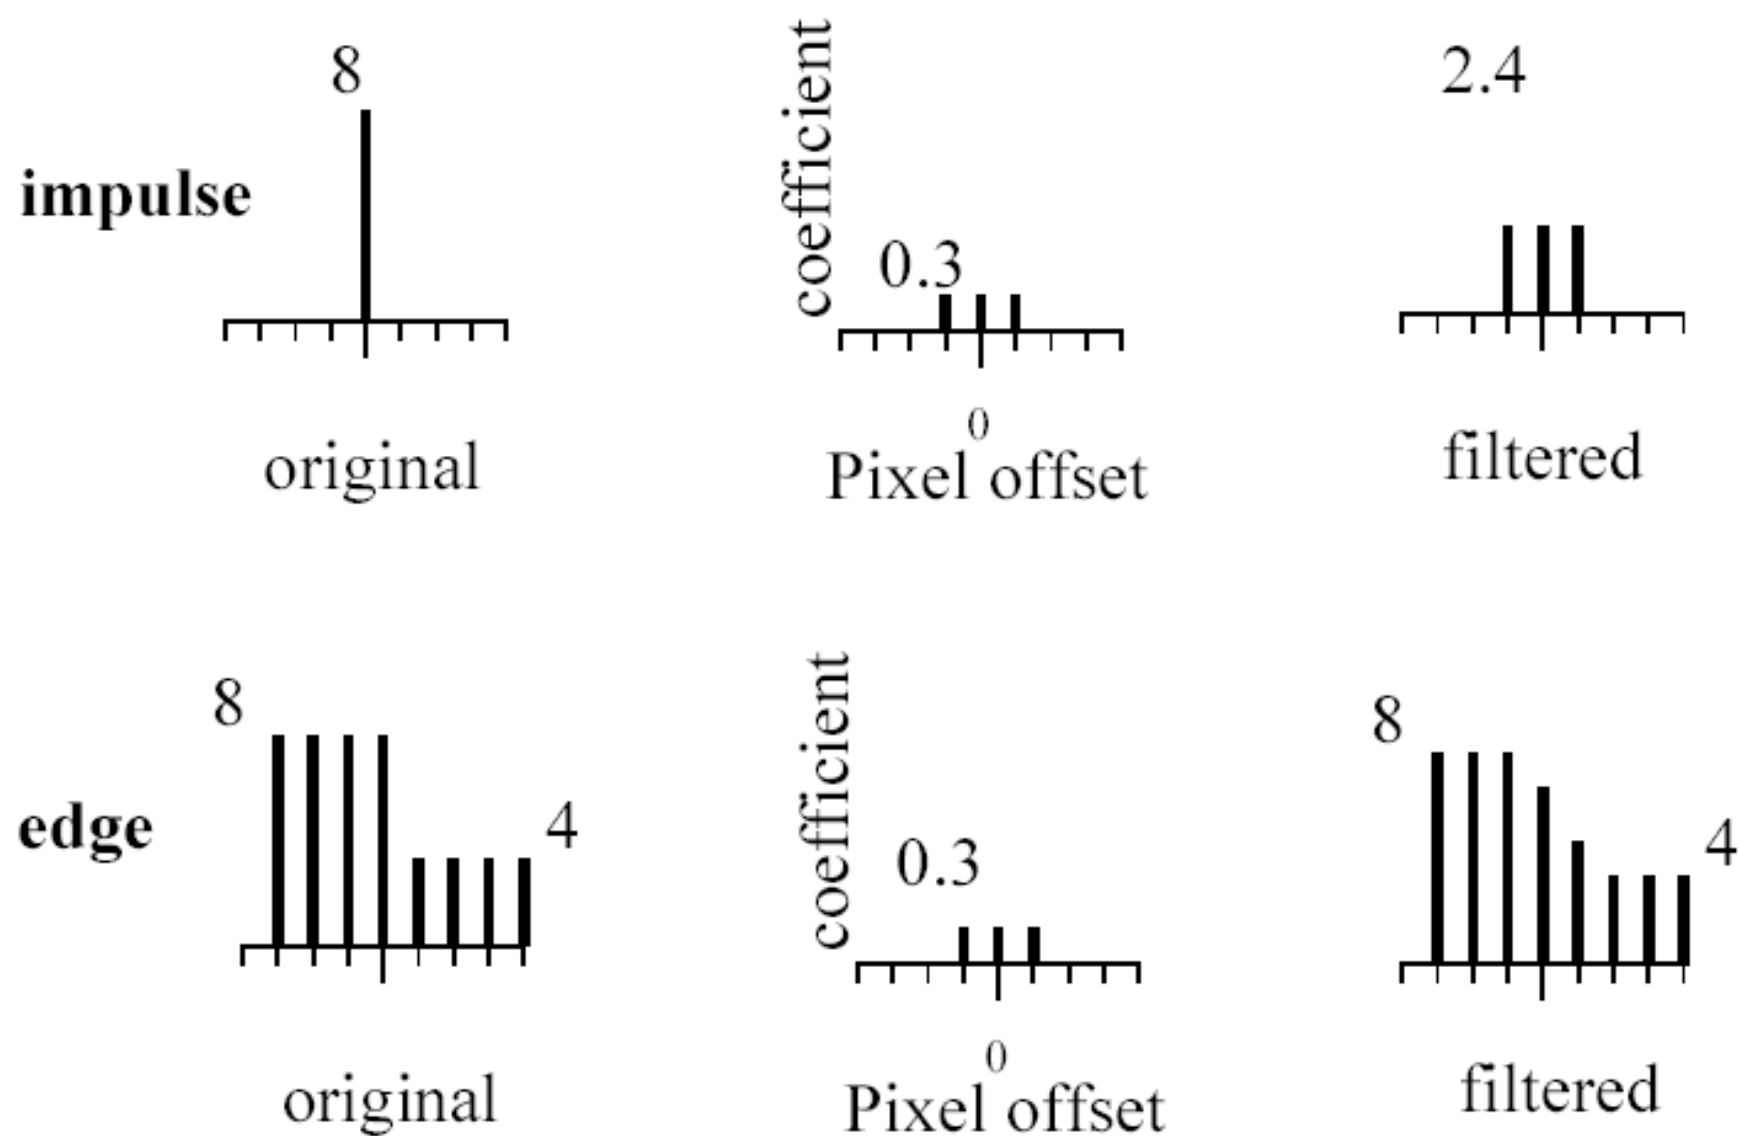
Blurred (filter applied in both dimensions).




# Blur examples




# Blur examples




# Linear filtering (warm-up slide)



original




—




?





# Linear filtering (no change)



original




—




Filtered  
(no change)


# Linear filtering



original




—





?

(remember blurring)




original





Blurred (filter applied in both dimensions).



# Sharpening





original




Sharpened  
original

# Sharpening example




Sharpened  
(differences are  
accentuated; constant  
areas are left untouched).

# Sharpening



before



after



# Simple 3x3 box blur

```
float input[(WIDTH+2) * (HEIGHT+2)];
```



```
float output[WIDTH * HEIGHT];
```

```
float weights[] = {1./9, 1./9, 1./9,  
                  1./9, 1./9, 1./9,  
                  1./9, 1./9, 1./9};
```

```
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int jj=0; jj<3; jj++)  
            for (int ii=0; ii<3; ii++)  
                tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];  
        output[j*WIDTH + i] = tmp;  
    }  
}
```

Will ignore boundary pixels today and  
assume output image is smaller than  
input (makes convolution loop bounds  
much simpler to write)

# 7x7 box blur



# Gaussian blur



- Obtain filter coefficients from sampling 2D Gaussian

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

- Produces weighted sum of neighboring pixels (contribution falls off with distance)
  - Truncate filter beyond certain distance

$$\begin{bmatrix} .075 & .124 & .075 \\ .124 & .204 & .124 \\ .075 & .124 & .075 \end{bmatrix}$$

# 7\*7 Gaussian blur



**7x7 box blur**







# What does convolution with this filter do?

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Sharpens image!

# 3x3 sharpen filter

Original



Sharpened



# What does convolution with these filters do?

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$


Extracts horizontal  
gradients

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$


Extracts vertical  
gradients



# Gradient detection filters



Horizontal gradients



Vertical gradients

Note: you can think of a filter as a “detector” of a pattern, and the magnitude of a pixel in the output image as the “response” of the filter to the region surrounding each pixel in the input image (this is a common interpretation in computer vision)



# Cost of convolution with N x N filter?

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float output[WIDTH * HEIGHT];
```

In this 3x3 box blur example:

Total work per image = 9 x WIDTH x HEIGHT

```
float weights[] = {1./9, 1./9, 1./9,  
                  1./9, 1./9, 1./9,  
                  1./9, 1./9, 1./9};
```

For N x N filter:  $N^2$  x WIDTH x HEIGHT

```
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int jj=0; jj<3; jj++)  
            for (int ii=0; ii<3; ii++)  
                tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];  
        output[j*WIDTH + i] = tmp;  
    }  
}
```

# Separable filter

- A filter is separable if it is the product of two other filters
- Example: a 2D box blur

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix}$$

- Exercise: write 2D gaussian and vertical/horizontal gradient detection filters as product of 1D filters (they are separable!)
- Key property: 2D convolution with separable filter can be written as two 1D convolutions!

# Implementation of 2D box blur via two 1D convolutions

```
int WIDTH = 1024
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1./3, 1./3, 1./3};

for (int j=0; j<(HEIGHT+2); j++)
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int ii=0; ii<3; ii++)
            tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];
        tmp_buf[j*WIDTH + i] = tmp;
    }

for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];
        output[j*WIDTH + i] = tmp;
    }
}
```

Total work per image =  $6 \times \text{WIDTH} \times \text{HEIGHT}$

For  $N \times N$  filter:  $2N \times \text{WIDTH} \times \text{HEIGHT}$

Extra cost of this approach?

Storage!

Challenge: can you achieve this work complexity without incurring this cost?

# Data-dependent filter (not a convolution)

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float output[WIDTH * HEIGHT];  
  
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float min_value = min( min(input[(j-1)*WIDTH + i], input[(j+1)*WIDTH + i]),  
                                min(input[j*WIDTH + i-1], input[j*WIDTH + i+1]) );  
        float max_value = max( max(input[(j-1)*WIDTH + i], input[(j+1)*WIDTH + i]),  
                                max(input[j*WIDTH + i-1], input[j*WIDTH + i+1]) );  
        output[j*WIDTH + i] = clamp(min_value, max_value, input[j*WIDTH + i]);  
    }  
}
```

- This filter clamps pixels to the min/max of its cardinal neighbors (e.g., hot-pixel suppression)

# Median filter

- Replace pixel with median of its neighbors
  - Useful noise reduction filter: unlike Gaussian blur, one bright pixel doesn't drag up the average for entire region
- Not linear, not separable
  - Filter weights are 1 or 0 (depending on image content)

```
uint8 input[(WIDTH+2) * (HEIGHT+2)];
uint8 output[WIDTH * HEIGHT];
for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        output[j*WIDTH + i] =
            // compute median of pixels
            // in surrounding 5x5 pixel window
    }
}
```

- **Basic algorithm for NxN support region:**
  - Sort  $N^2$  elements in support region, pick median  $O(N^2 \log(N^2))$  work per pixel



original image



1px median filter



3px median filter



10px median filter



