

# **Computer Graphics**

## **- Bilateral & Nonlocal Mean Filtering**

Junjie Cao @ DLUT

Spring 2016

<http://jjcao.github.io/ComputerGraphics/>

# Bilateral filtering



# Introduction

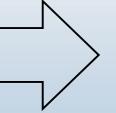
# Goal: Image Smoothing

Split an image into:

- large-scale features, structure
- small-scale features, texture



*input*



**BLUR**



*smoothed  
(structure, large scale)*

**HALOS**



*residual  
(texture, small scale)*

**Gaussian Convolution**

# Impact of Blur and Halos

- If the decomposition introduces blur and halos, the final result is corrupted.

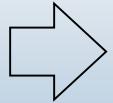
Sample manipulation:  
increasing texture  
(residual  $\times 3$ )



# Bilateral Filter (no Blur, no Halos) vs Naïve Approach: Gaussian Blur (with blur & halos)



*input*



*smoothed*  
*(structure, large scale)*

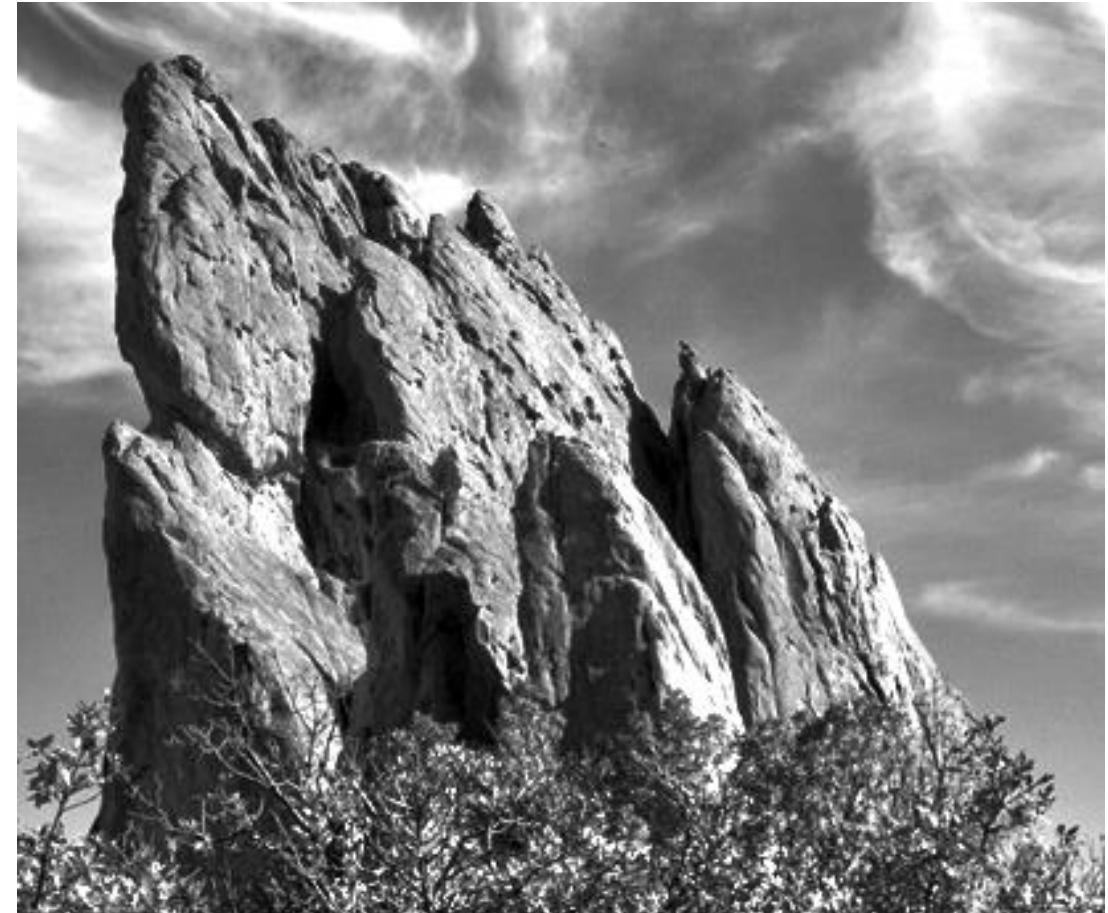


*residual*  
*(texture, small scale)*

**edge-preserving: Bilateral Filter**



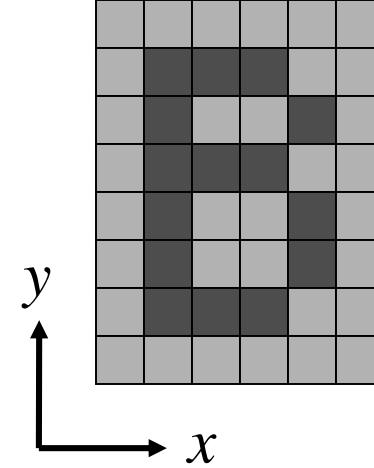
increasing texture  
with Gaussian convolution  
**H A L O S**



increasing texture  
with bilateral filter  
**N O H A L O S**

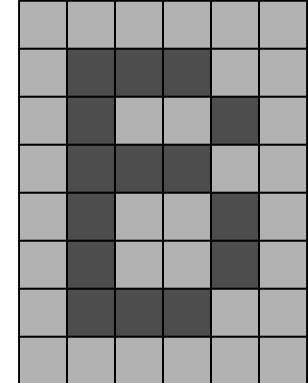
# Notation and Definitions

- Image = 2D array of pixels
- Pixel = intensity (scalar) or color (3D vector)
- $I_p$  = value of image  $I$  at position:  $\mathbf{p} = (p_x, p_y)$
- $F [ I ]$  = output of filter  $F$  applied to image  $I$



# What is filtering?

- Images are not smooth because adjacent pixels are different.
- Smoothing = making adjacent pixels look more similar.
- Smoothing strategy
  - pixel → average of its neighbors



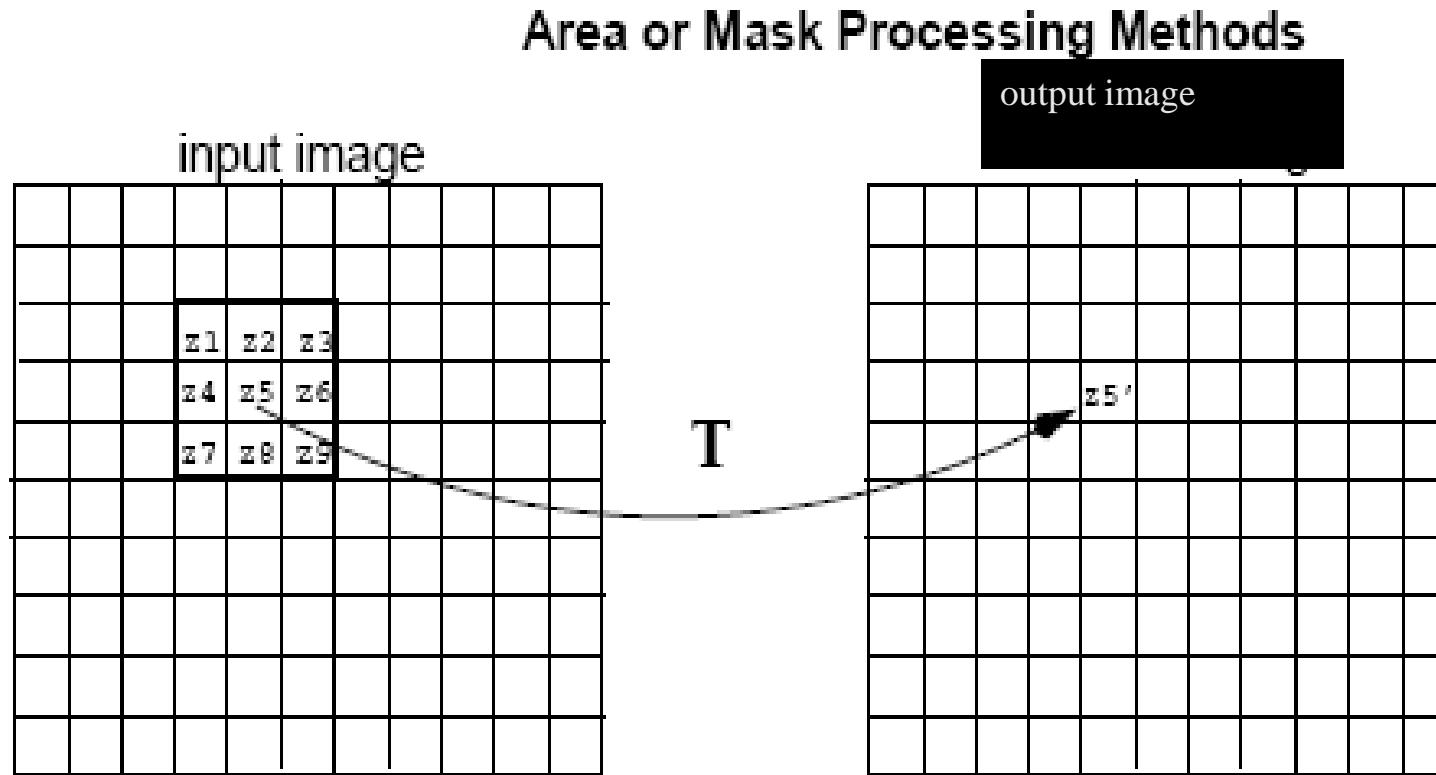
# Spatial filtering/Mask processing - Operation

Example: weighted sum of input pixels.

$$N5' = R = w_1z_1 + w_2z_2 + \dots + z_9w_9$$

w1	w2	w3
w4	w5	w6
w7	w8	w9

mask  
weights:



$$g(x,y) = T[f(x,y)]$$

T operates on a neighborhood of pixels

A **filtered image** is generated as the **center** of the mask moves to every pixel in the input image.

# Why is it important?

- Many applications with high quality results.
  - Computational photography
  - Computer graphics
- Papers in top conferences
  - Siggraph 15: An L1 Image Transform for Edge-Preserving Smoothing and Scene-Level Intrinsic Decomposition
  - Siggraph asia 15: Rolling Guidance Normal Filter for Geometric Processing
  - Siggraph 14: Bilateral texture filtering
  - Siggraph 14: Fast Local Laplacian Filters: Theory and Applications
  - Siggraph asia 14: Depth of field rendering via adaptive recursive filtering
  - ...

# Photographic Style Transfer [Bae 06]



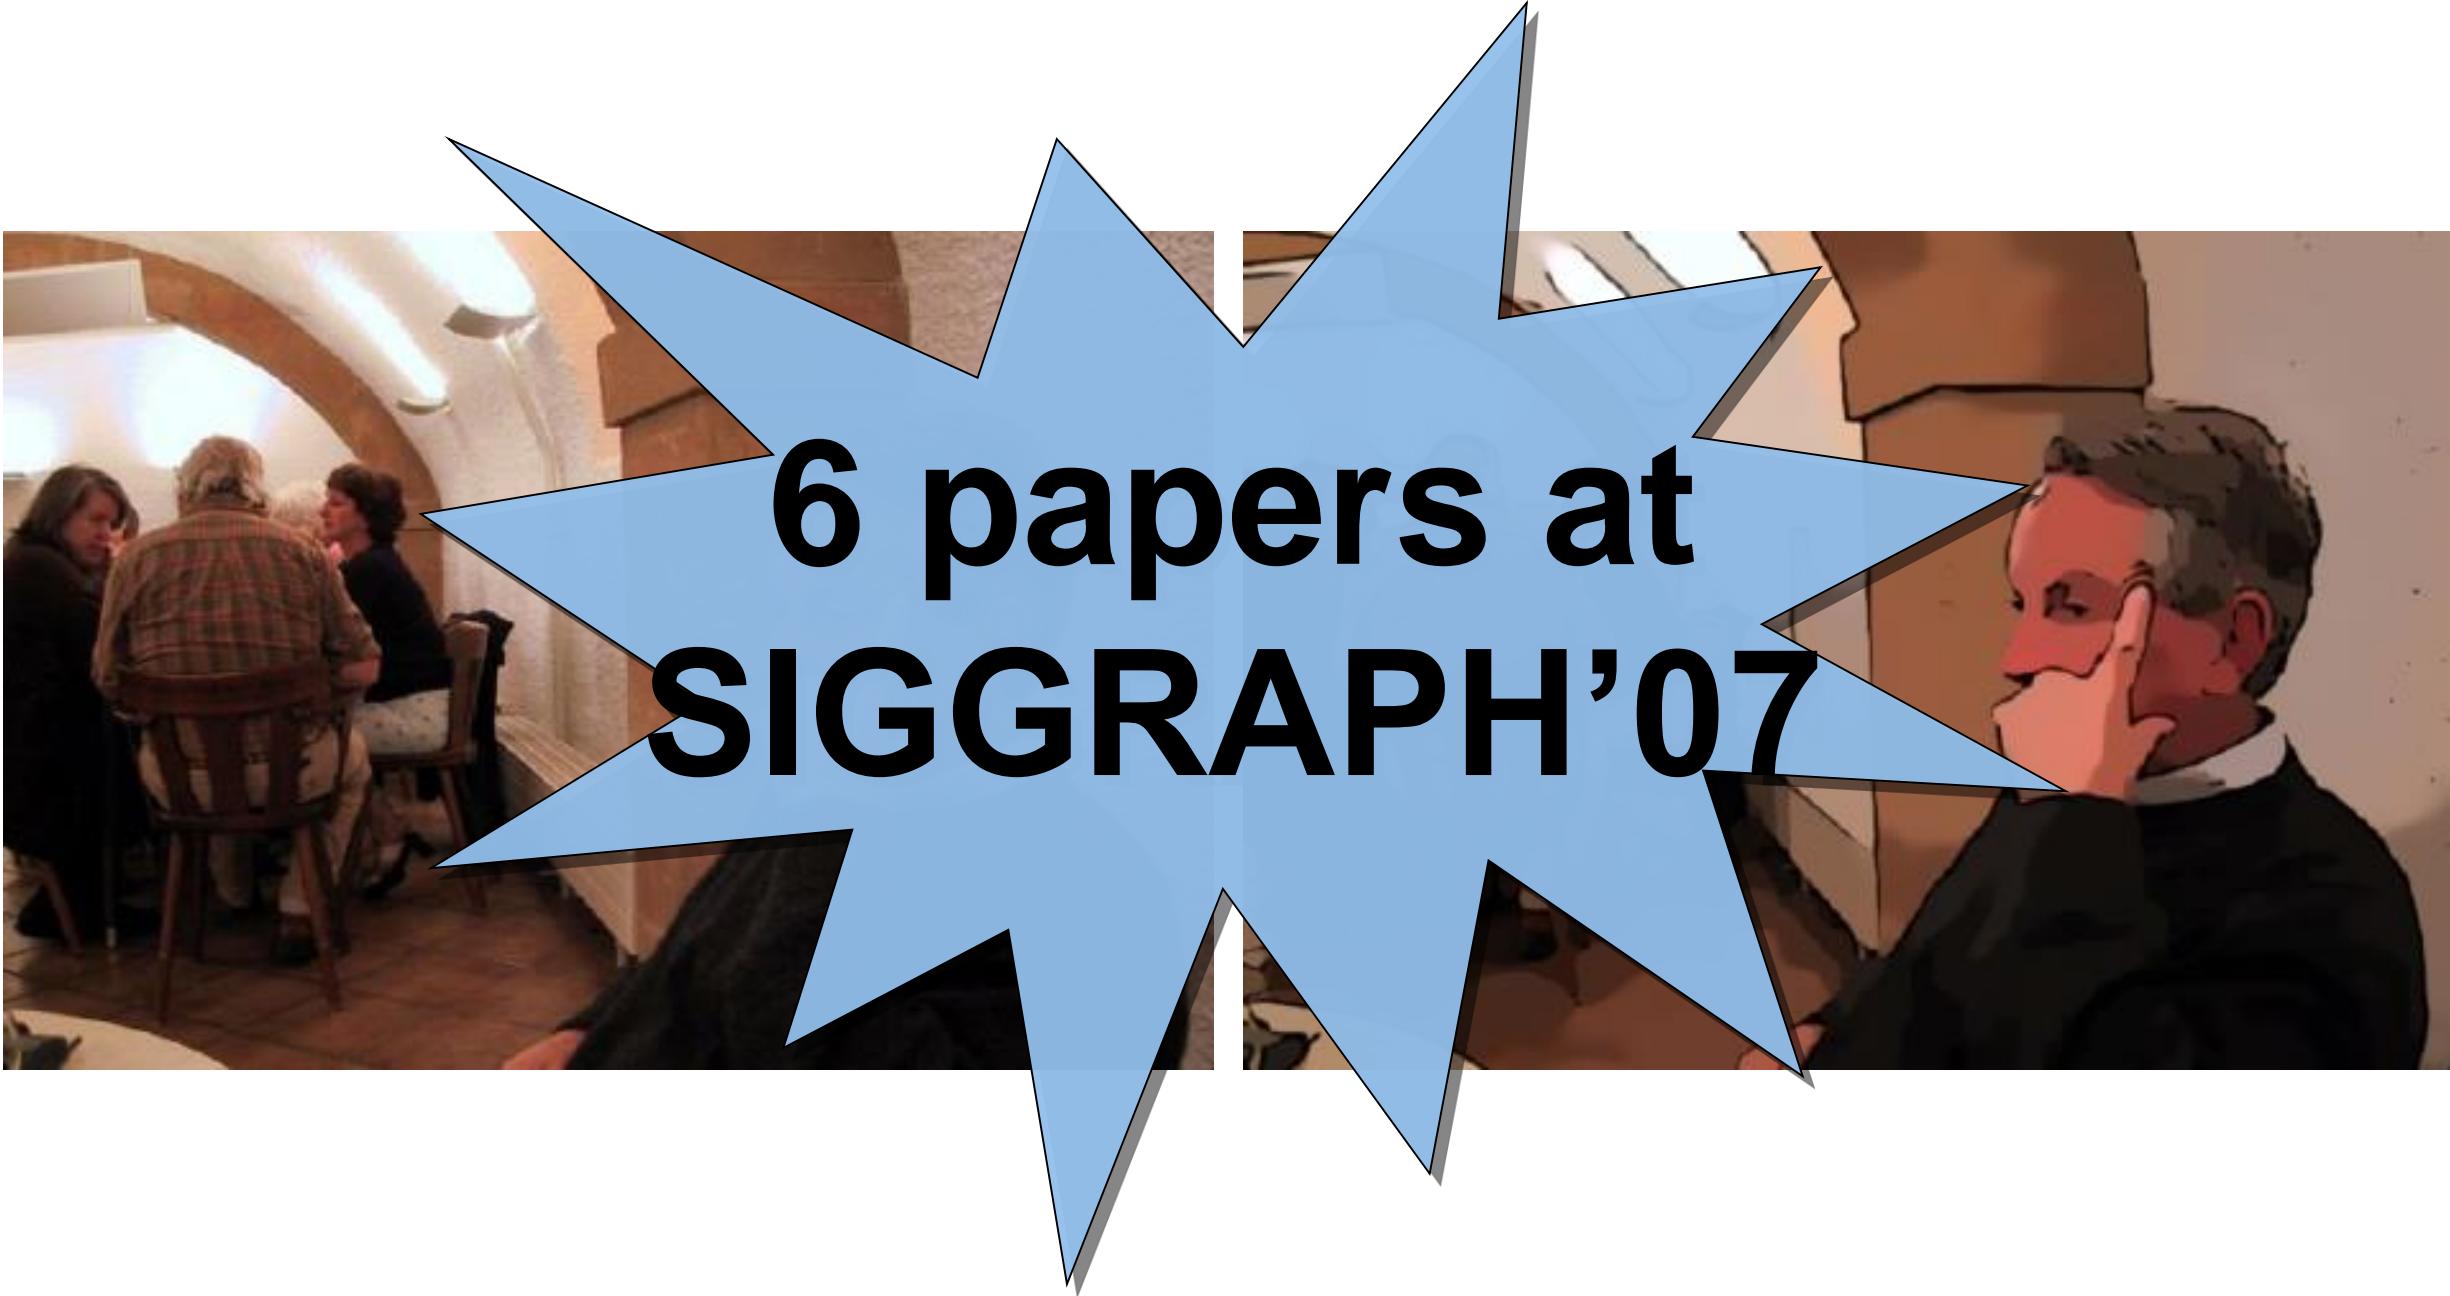
# Tone Mapping [Durand 02]



HDR input



# Cartoon Rendition [Winnemöller 06]



# Intrinsic image decomposition



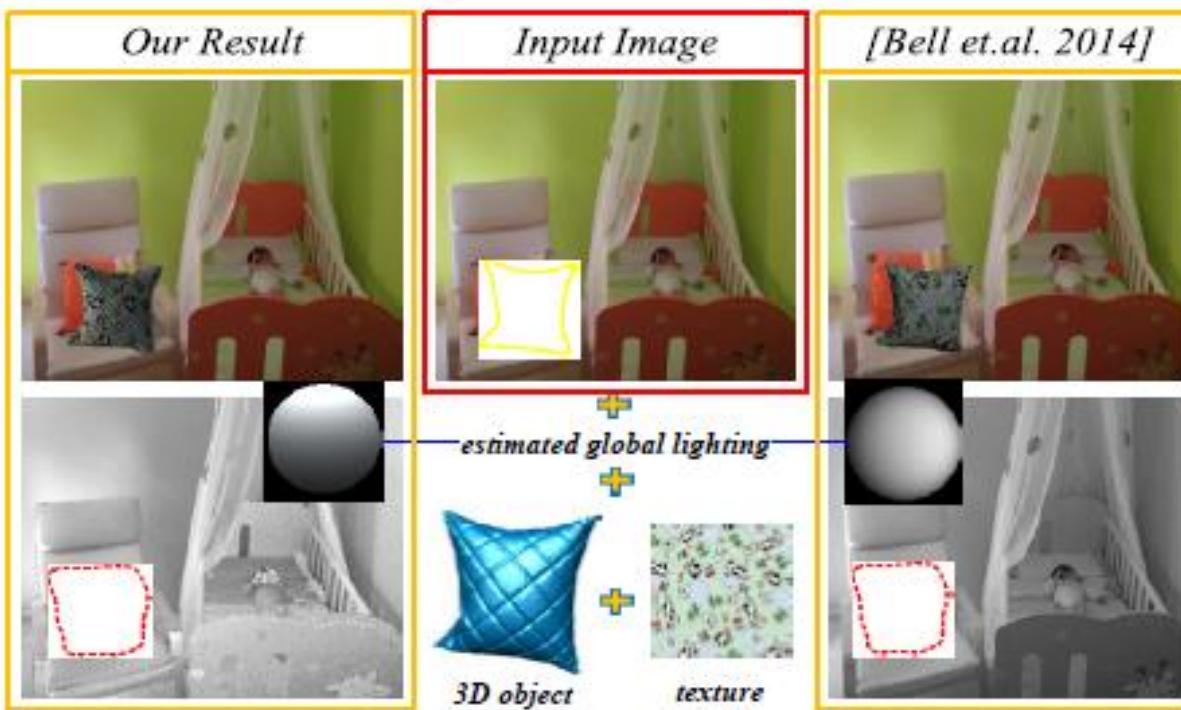
(e) *Original Image*

(f) *Image Flattening*

(g) *Reflectance*

(h) *Shading*

(i) *Re-Texturing*



# Keywords & terms

- Smoothing, denoising, blur, decomposition
- Feature, Structure, texture
- Blur, halo
- Convolution, kernel
- Spatial Filtering,

# Many Other Options

- Bilateral filtering is not the only image smoothing filter
  - Diffusion, wavelets, Bayesian...
- We focus on bilateral filtering
  - Suitable for strong smoothing used in computational photography
  - Conceptually simple

# Content of the Course

All you need to know about bilateral filtering:

- Definition of the bilateral filter
- Parameter influence and settings
- Applications
- Relationship to other filters
- Theoretical properties
- Efficient implementation

Course webpage (google “bilateral filter course”):

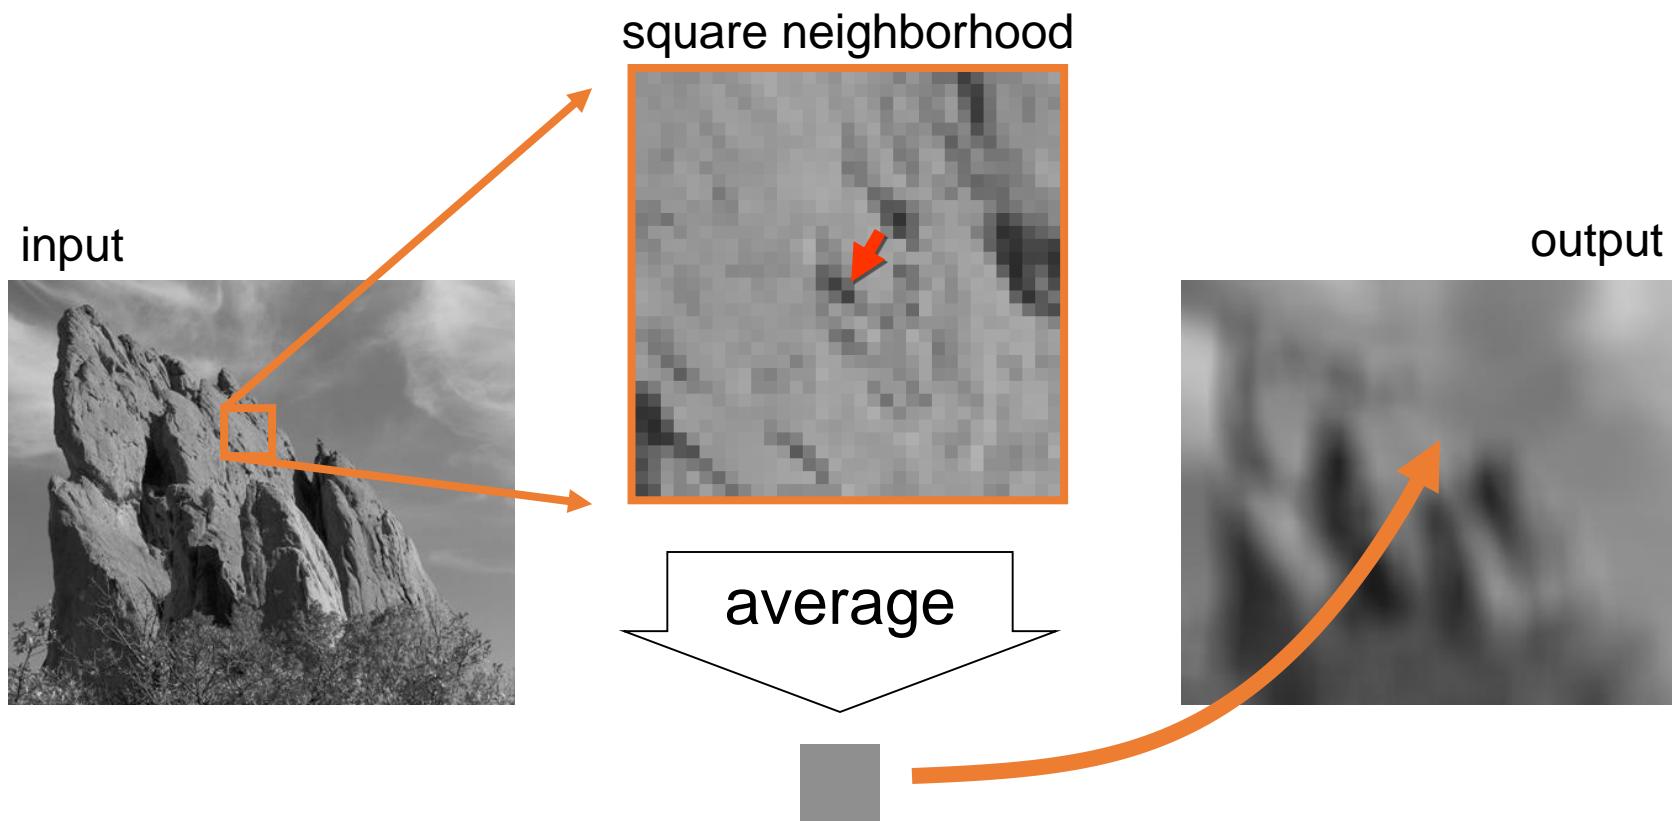
[http://people.csail.mit.edu/sparis/bf\\_course/](http://people.csail.mit.edu/sparis/bf_course/)

Detailed course notes

C++ and Matlab code

# Naïve Image Smoothing: Gaussian Blur

# Box Average



# Equation of Box Average

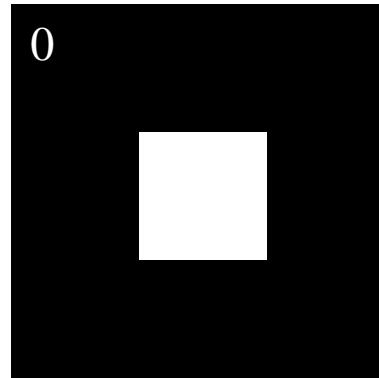
$$BA[I]_p = \sum_{q \in S} B_\sigma(p - q) I_q$$

result at pixel p

sum over all pixels q

intensity at pixel q

normalized box function

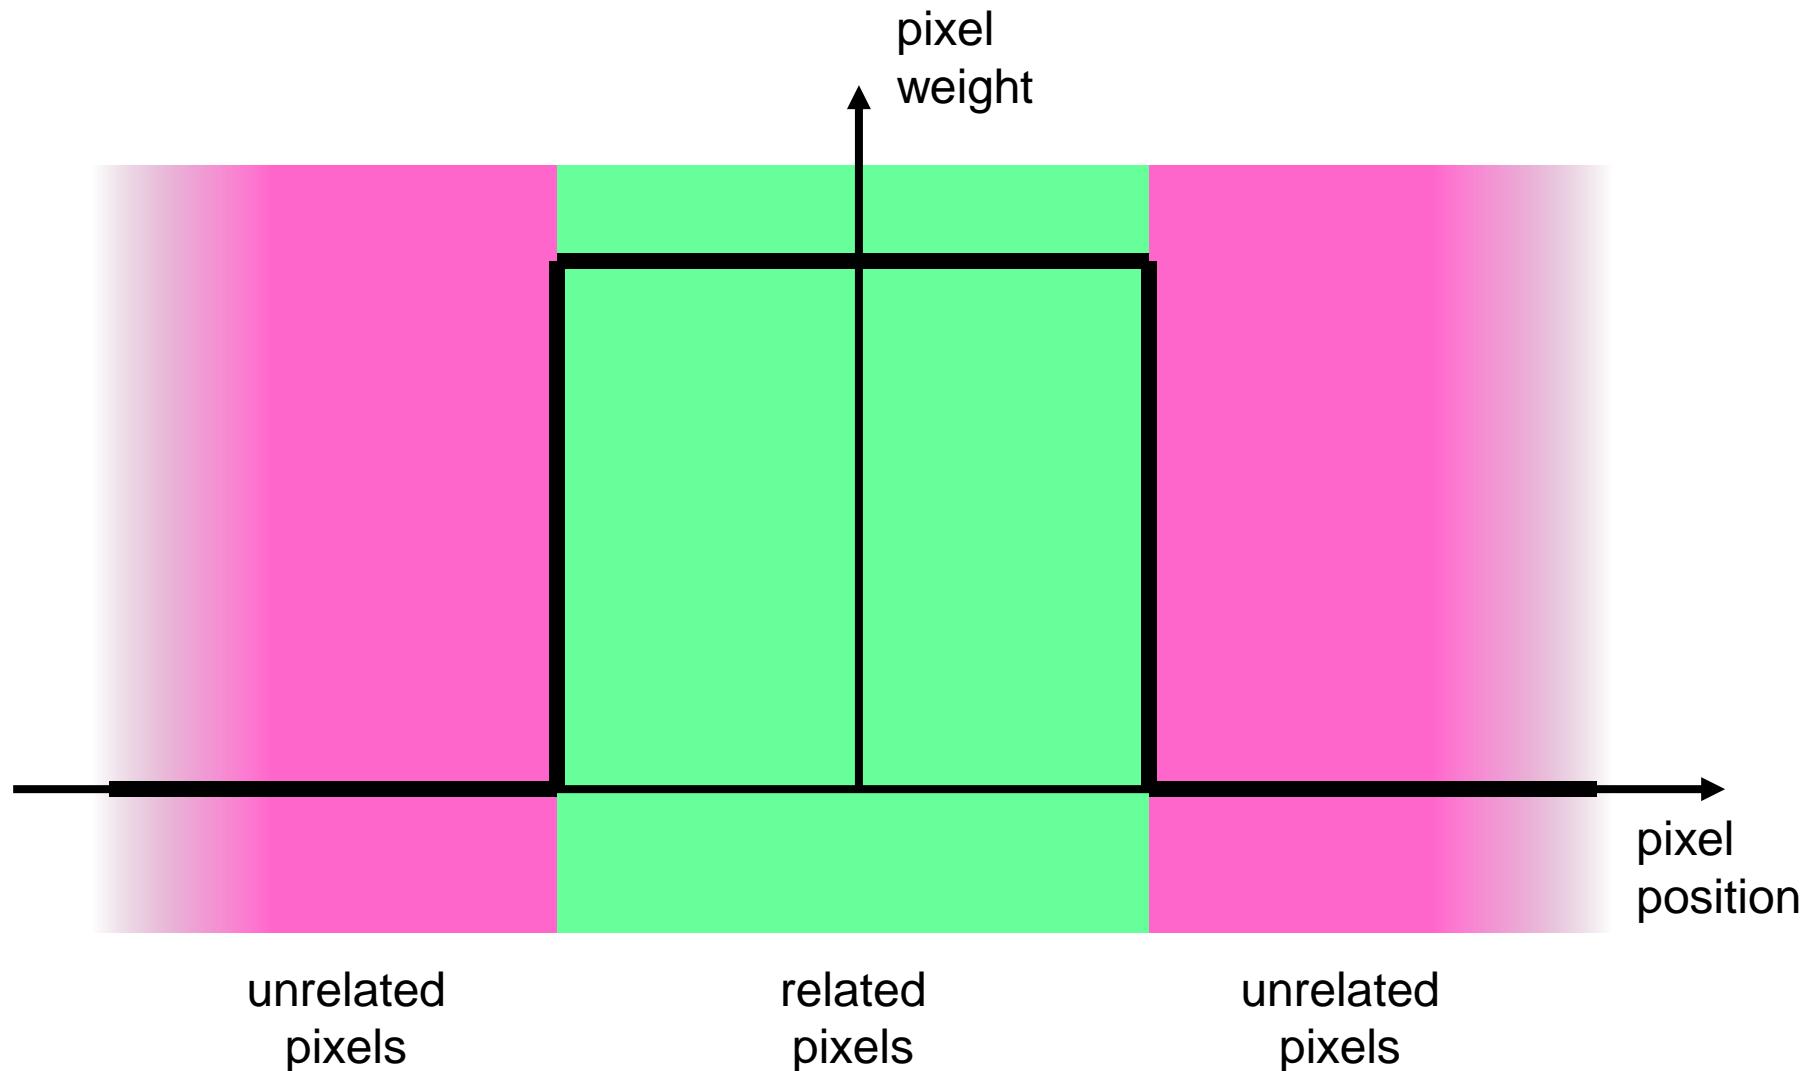


# Square Box Generates Defects

- Axis-aligned streaks
- Blocky results

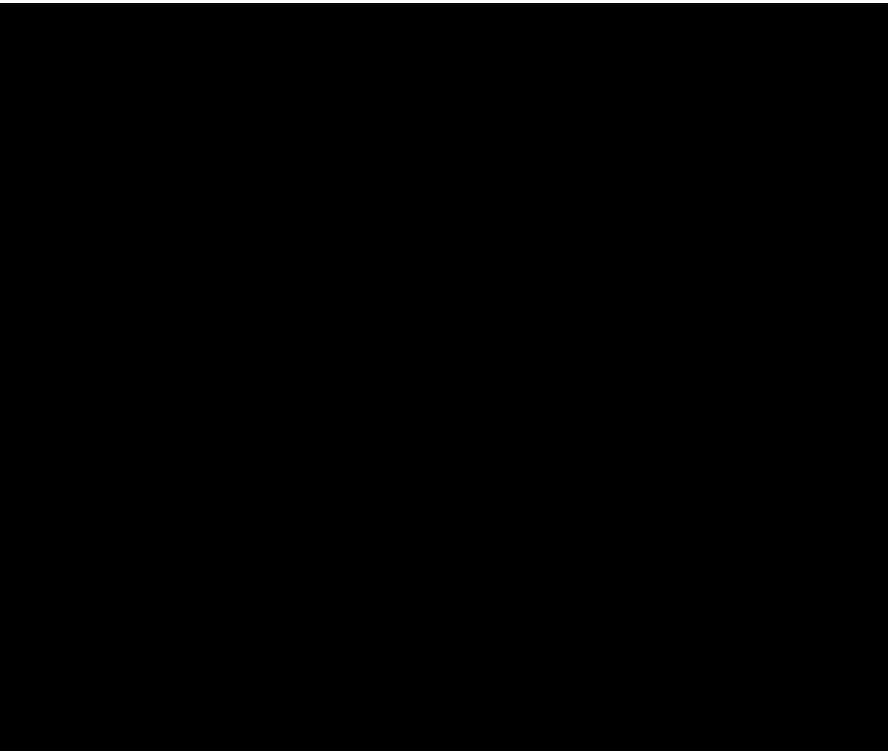


# Box Profile

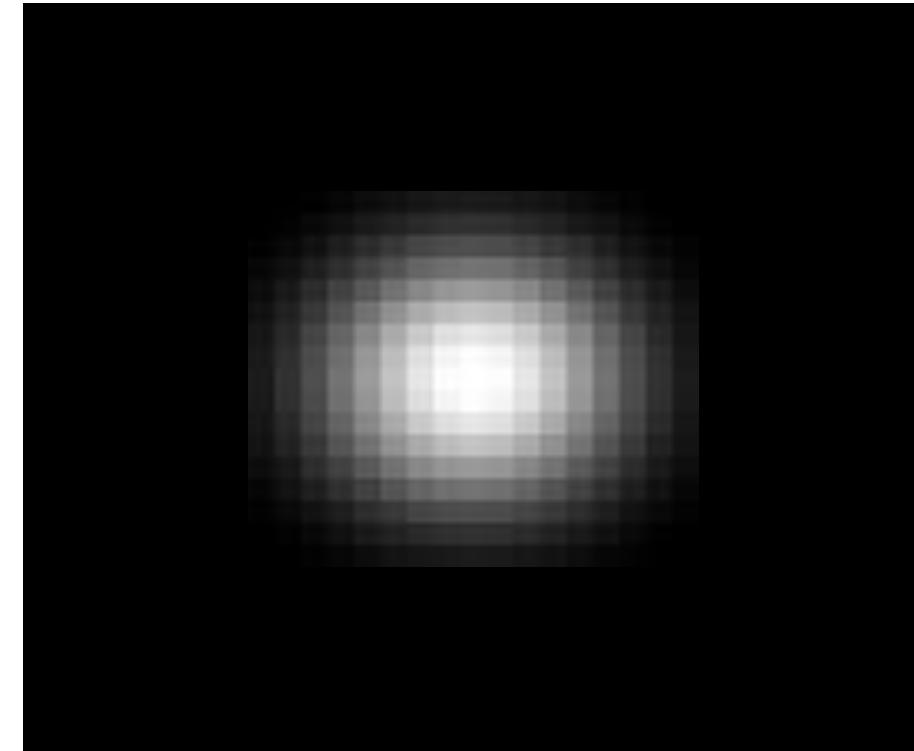


# Strategy to Solve these Problems

- Use an isotropic (*i.e.* circular) window.
- Use a window with a smooth falloff.

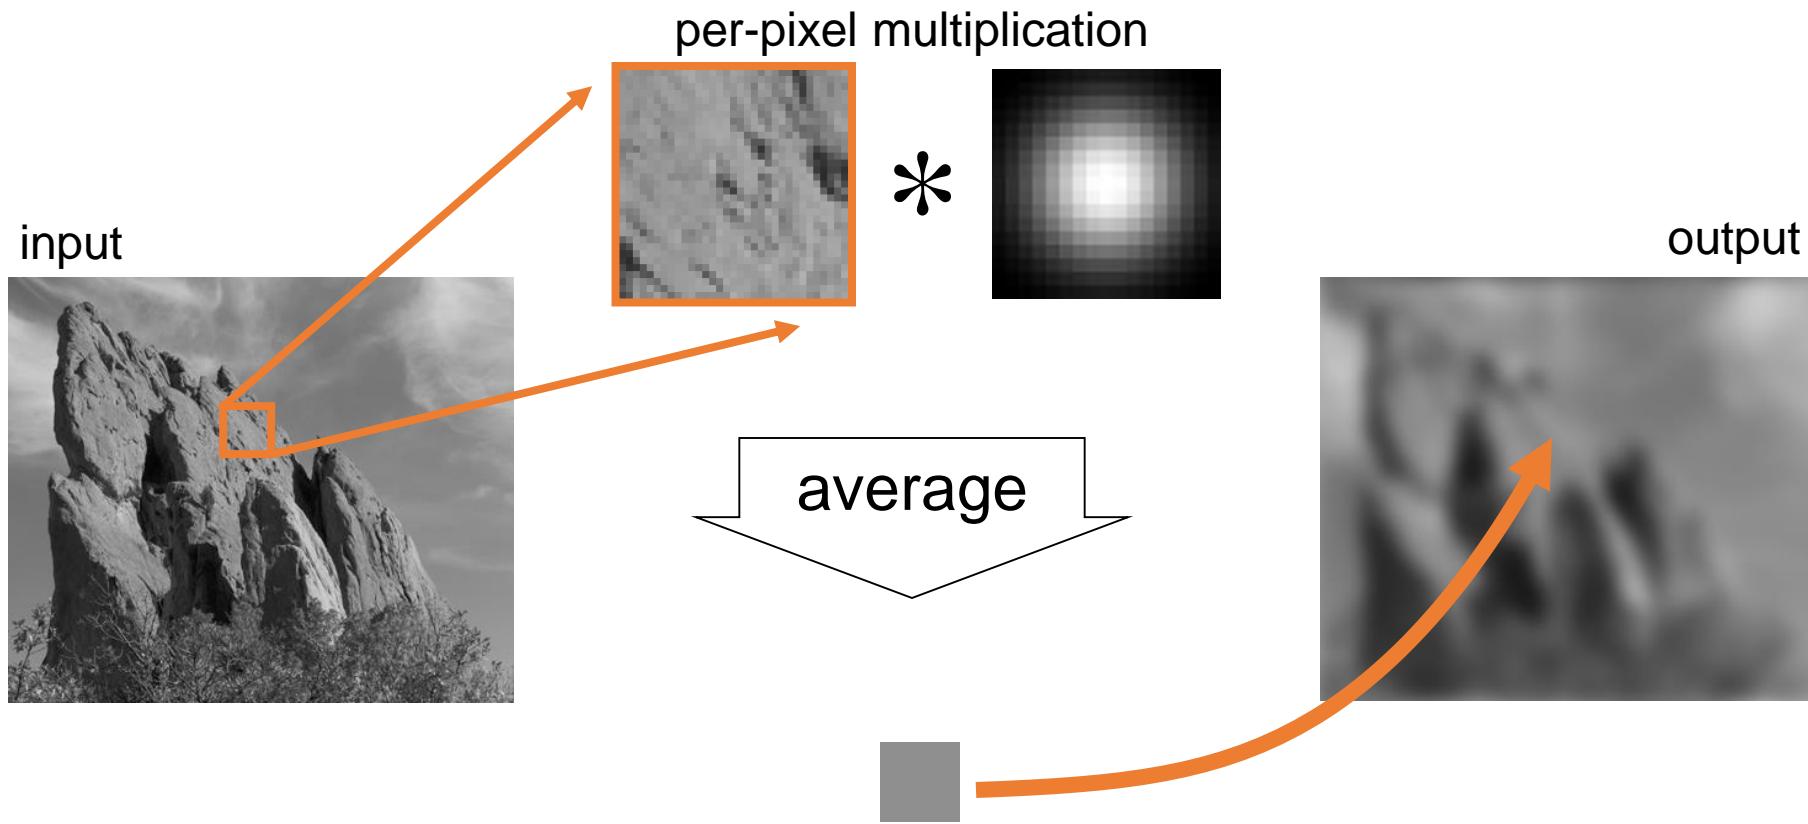


box window



Gaussian window

# Gaussian Blur



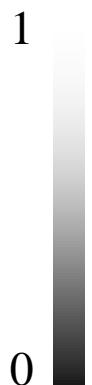
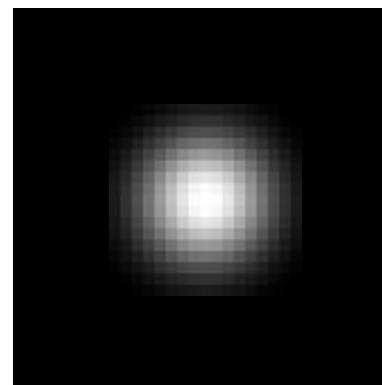
# Equation of Gaussian Blur

Same idea: **weighted average of pixels.**

$$GB [I]_p = \sum_{q \in S} G_\sigma(||p - q||) I_q$$

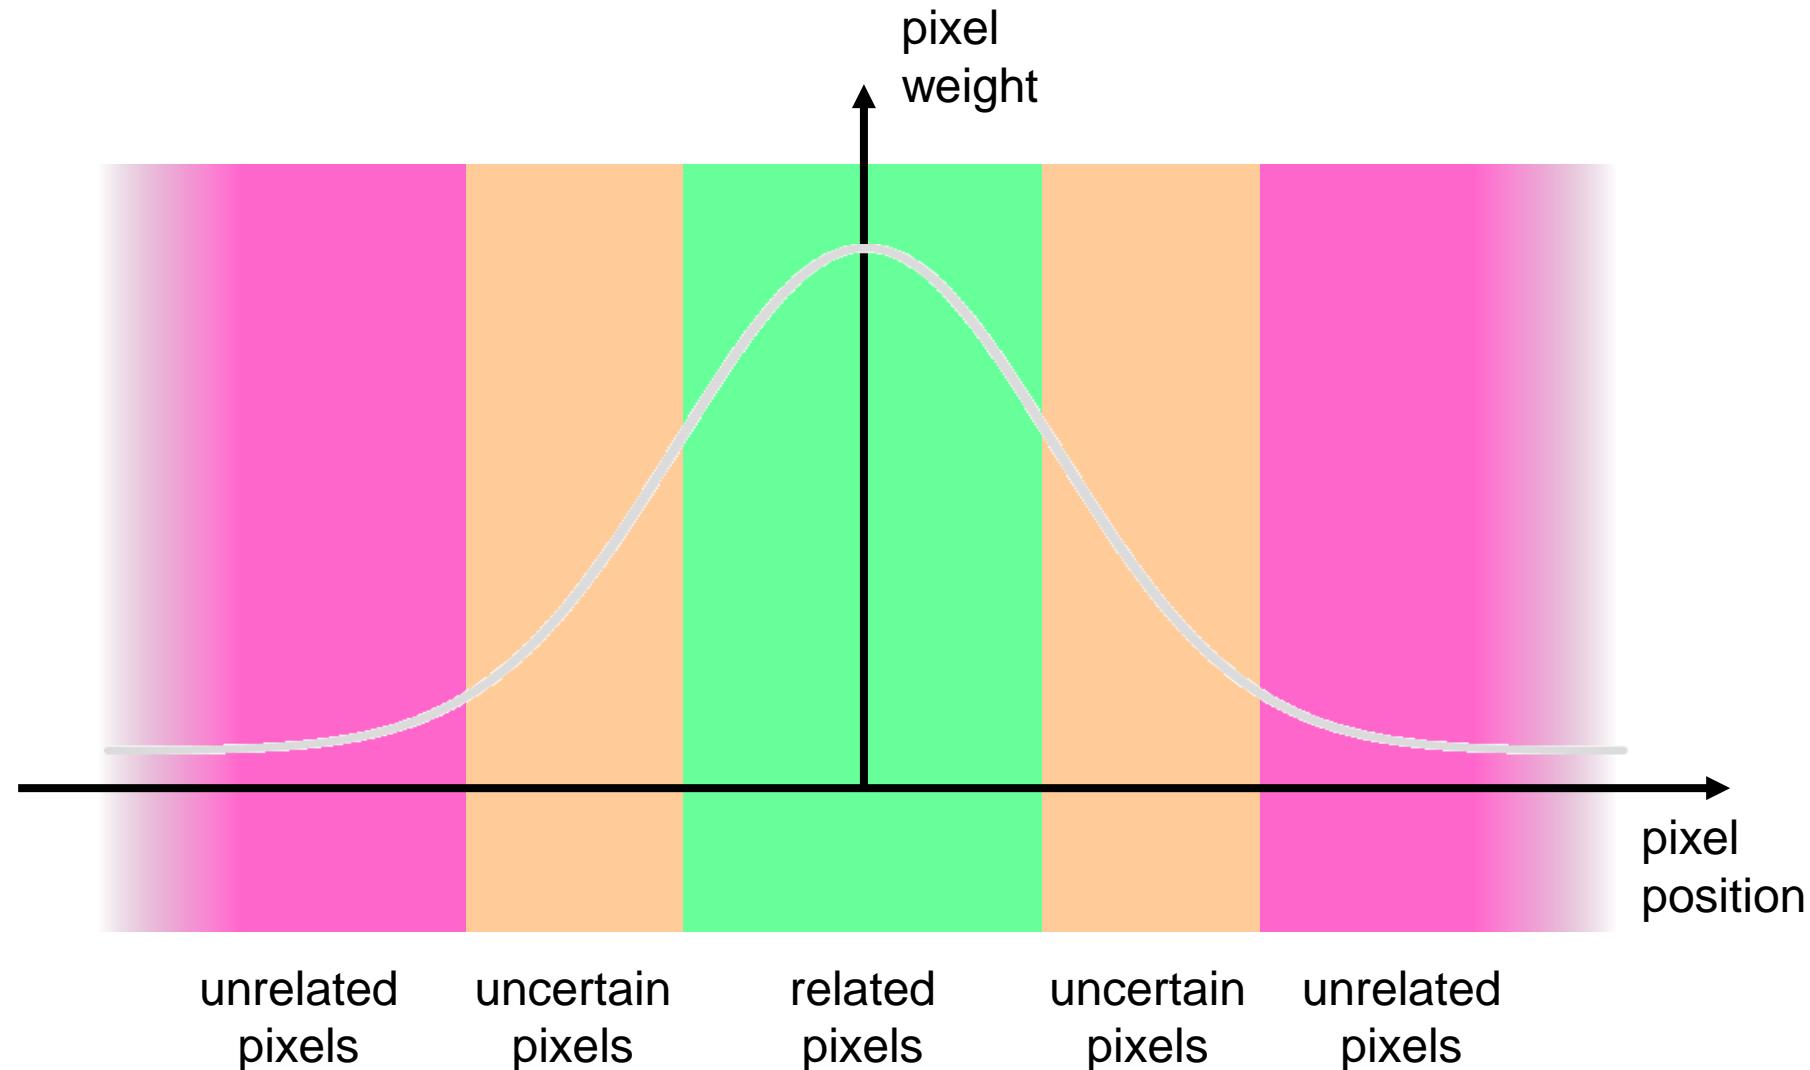


normalized  
Gaussian function



# Gaussian Profile

$$G_{\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{x^2}{2\sigma^2}\right)$$



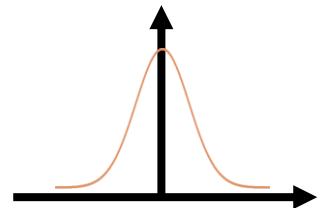
# Spatial Parameter



input

$$GB [I]_p = \sum_{q \in S} G_{\sigma}(| | p - q | |) I_q$$

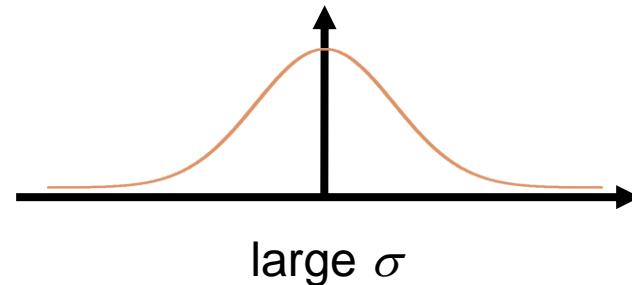
size of the window



small  $\sigma$



limited smoothing



large  $\sigma$



strong smoothing

# How to set $\sigma$

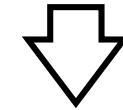
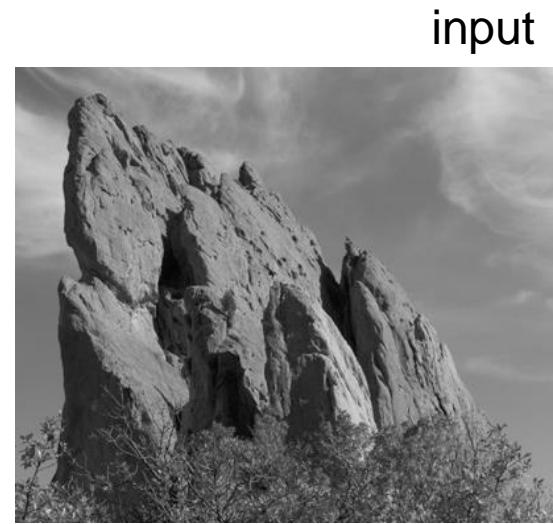
- Depends on the application.
- Common strategy: proportional to image size
  - e.g. 2% of the image diagonal
  - property: independent of image resolution

# Properties of Gaussian Blur

- Weights independent of spatial location
  - linear convolution
  - well-known operation
  - efficient computation (recursive algorithm, FFT...)

# Properties of Gaussian Blur

- Does smooth images
- But smoothes too much:  
**edges are blurred.**
  - Only spatial distance matters
  - No edge term



input  
output

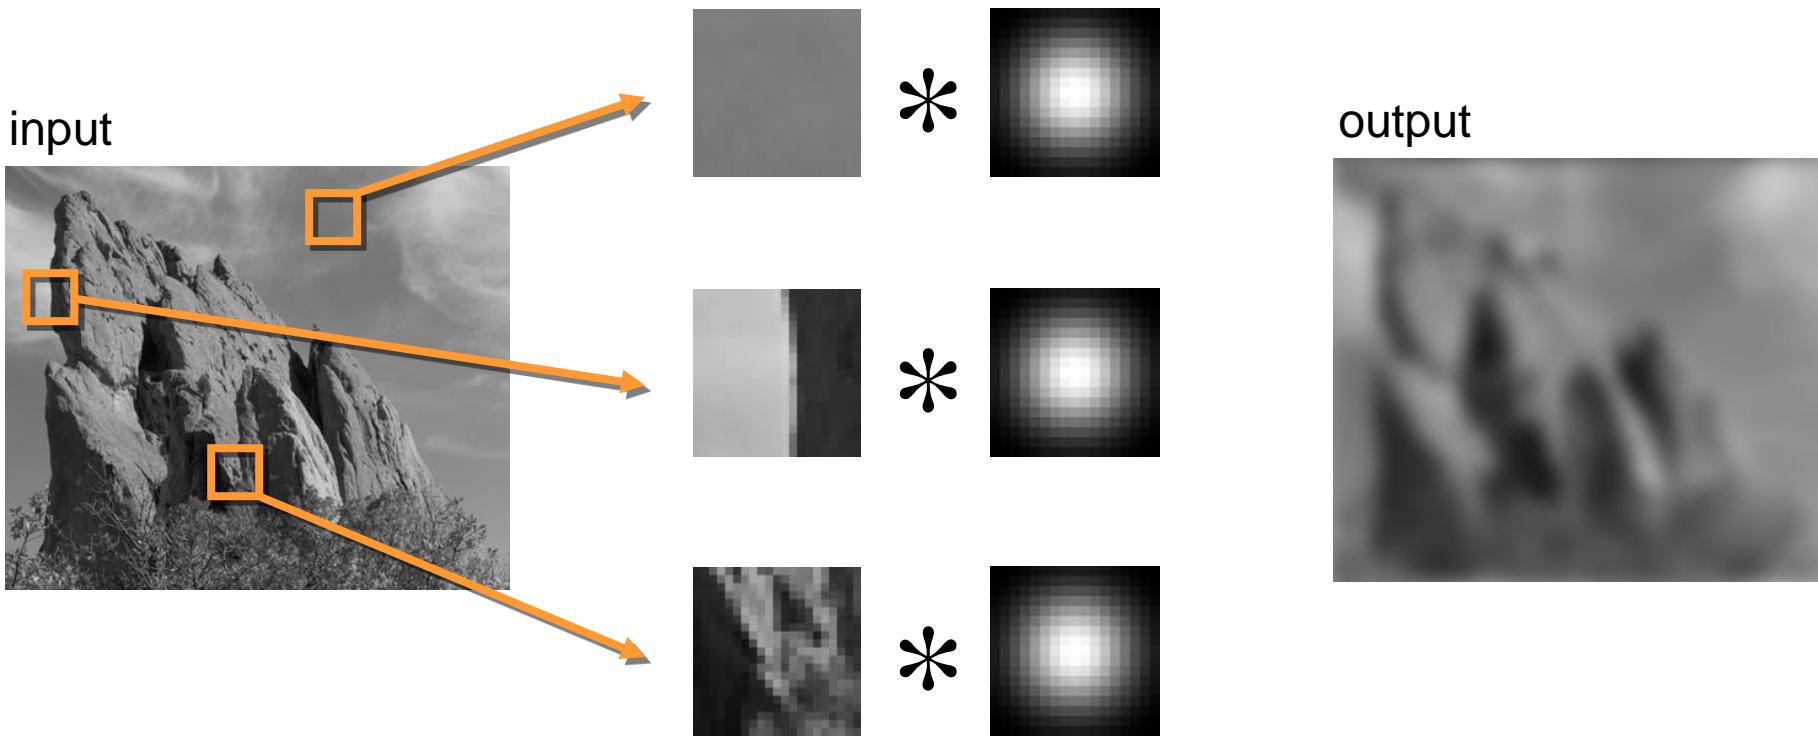


$$GB [I]_p = \sum_{q \in S} G_\sigma(||p - q||) I_q$$

space

# “Fixing the Gaussian Blur”: the Bilateral Filter

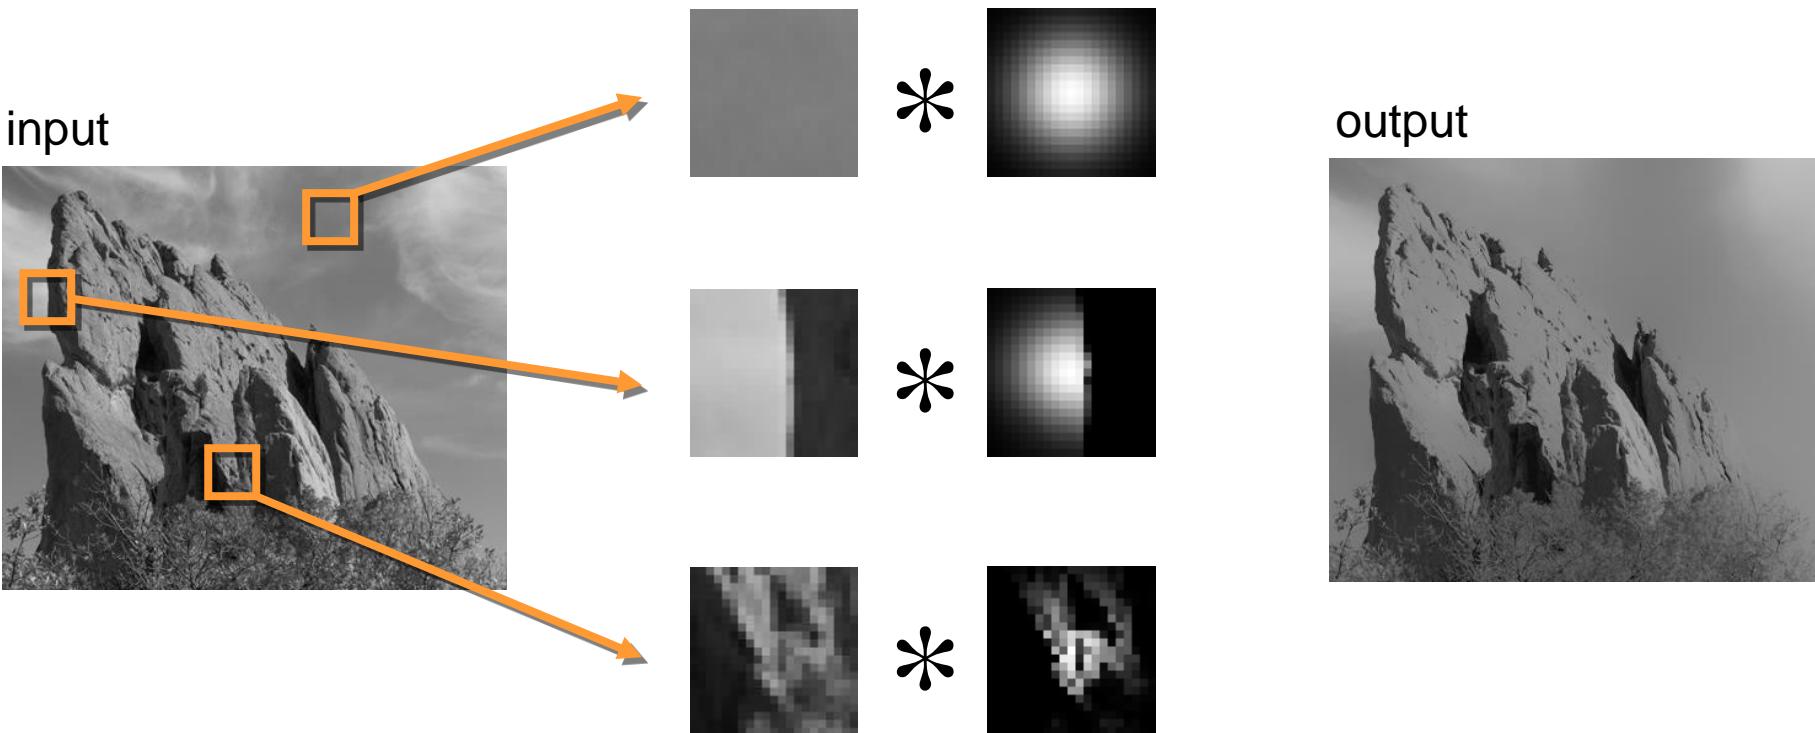
# Blur Comes from Averaging across Edges



Same Gaussian kernel everywhere.

# Bilateral Filter

## No Averaging across Edges



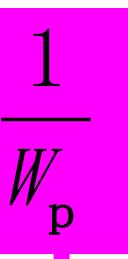
The kernel shape depends on the image content.

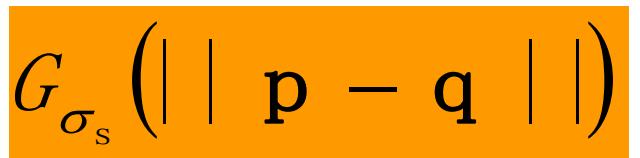
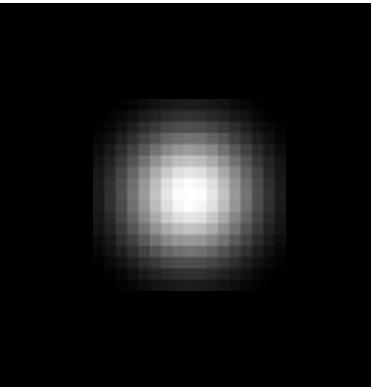
[Aurich 95, Smith 97, Tomasi 98]

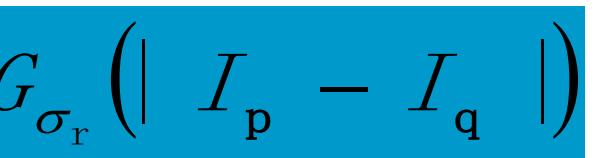
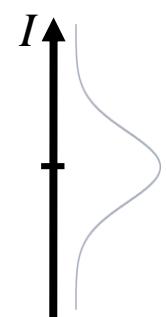
# Bilateral Filter Definition: an Additional Edge Term

Same idea: **weighted average of pixels.**

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(| | p - q | |) G_{\sigma_r}(| I_p - I_q |) I_q$$

new  
  
normalization factor

not new  
  
space weight  


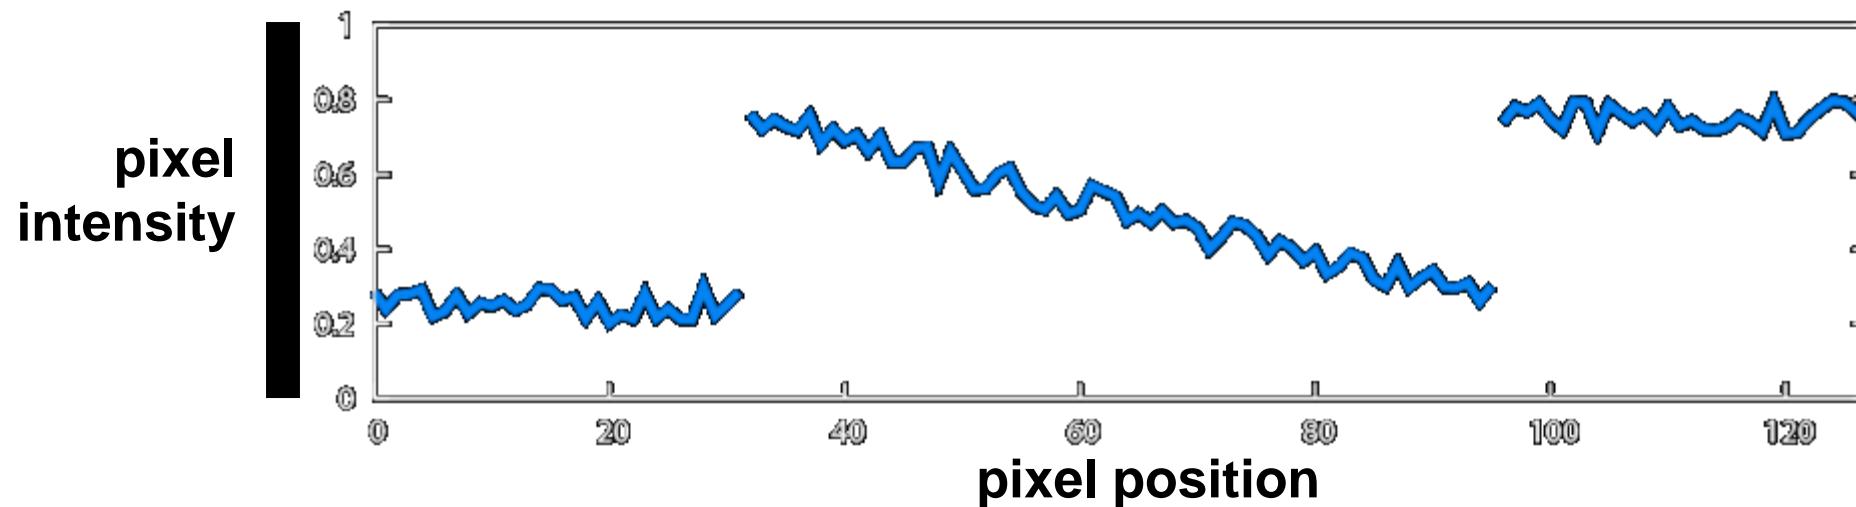
new  
  
range weight  


# Illustration a 1D Image

- 1D image = line of pixels

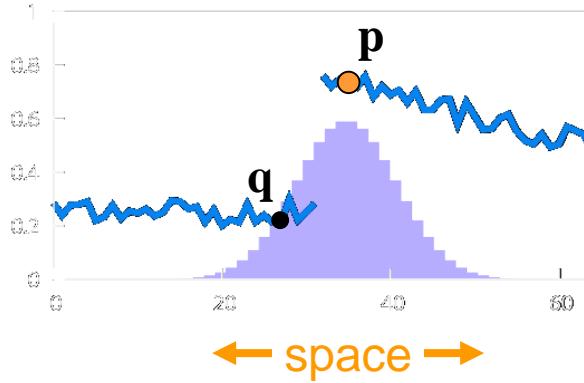


- Better visualized as a plot



# Gaussian Blur and Bilateral Filter

## Gaussian blur

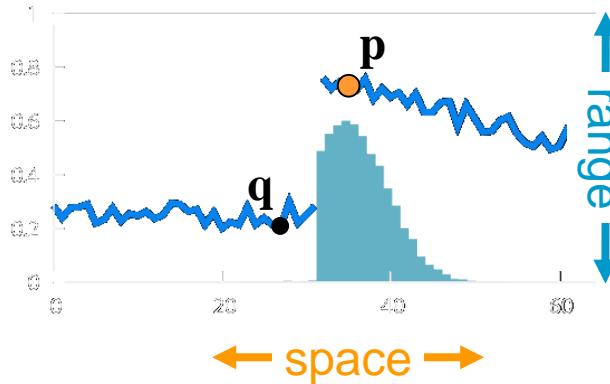


$$GB [I]_p = \sum_{q \in S} G_\sigma(||p - q||) I_q$$

space

## Bilateral filter

[Aurich 95, Smith 97, Tomasi 98]

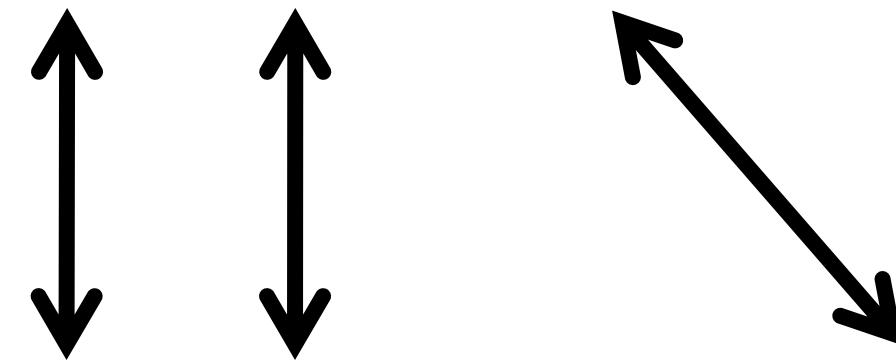


$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(||p - q||) G_{\sigma_r}(||I_p - I_q||) I_q$$

space

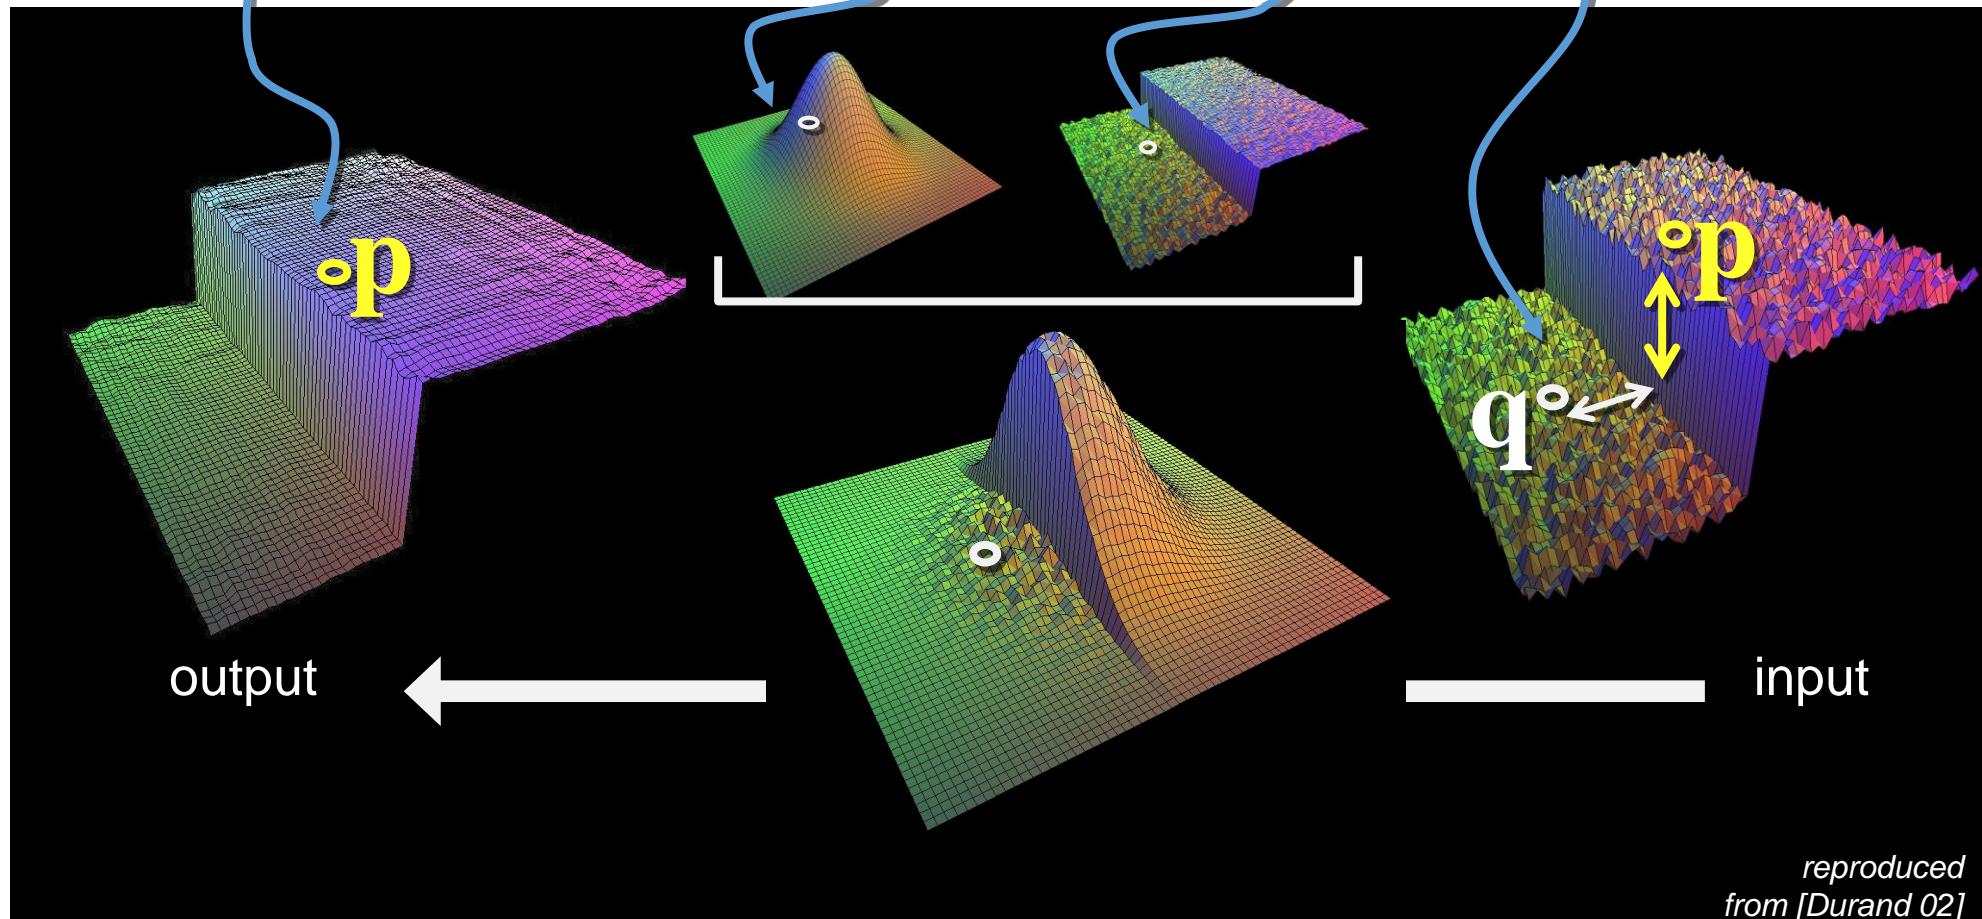
range

normalization



# Bilateral Filter on a Height Field

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|I_p - I_q|) I_q$$



reproduced  
from [Durand 02]

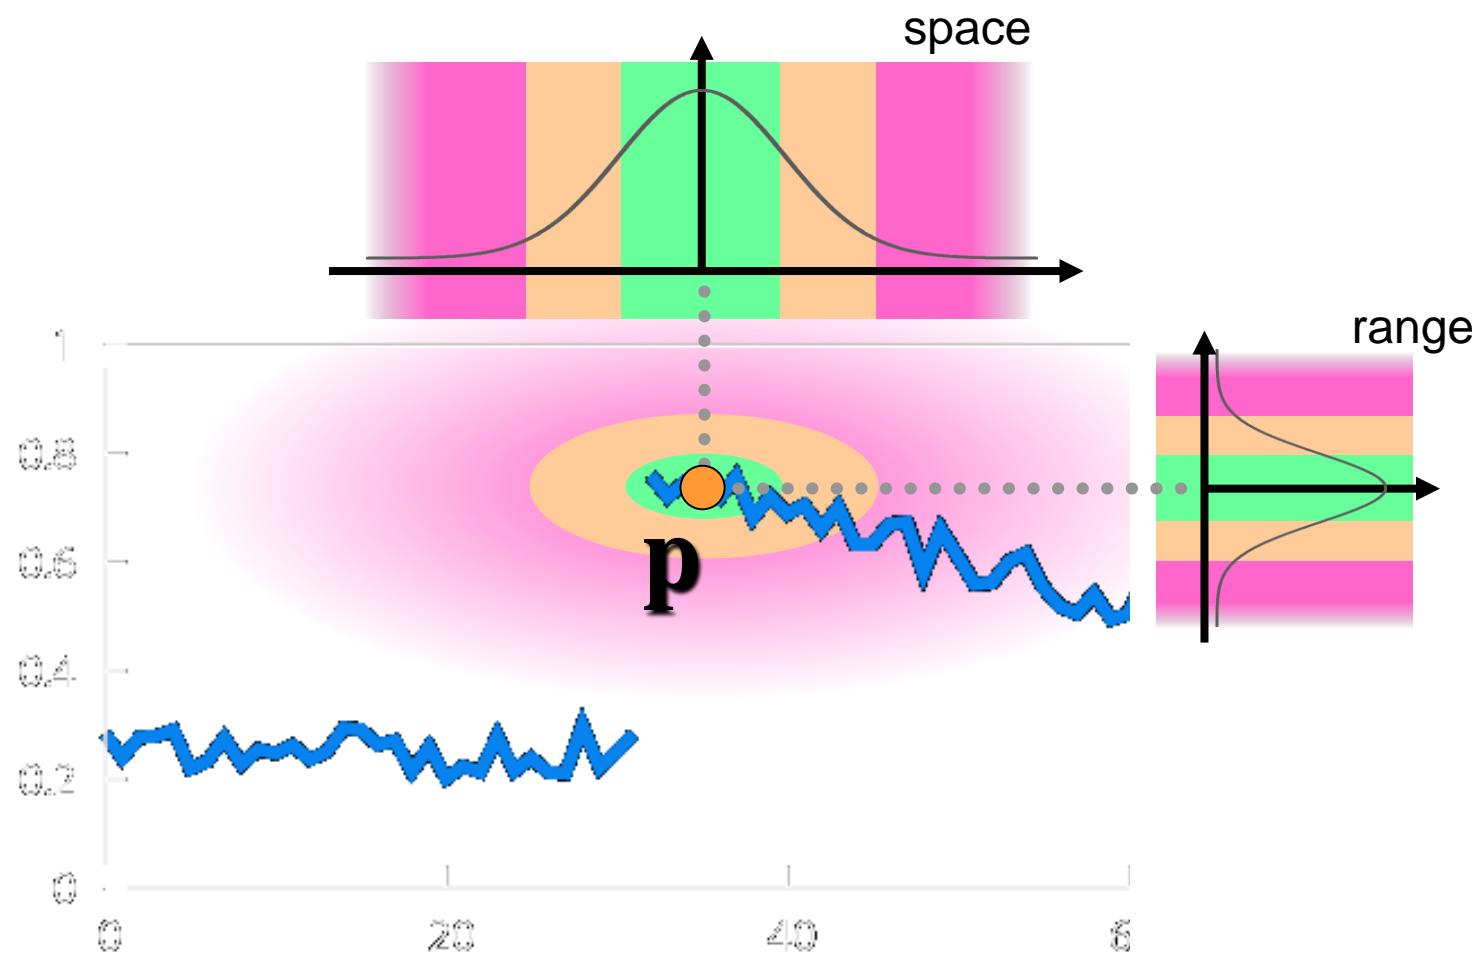
# Space and Range Parameters

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|I_p - I_q|) I_q$$


- space  $\sigma_s$ : spatial extent of the kernel, size of the considered neighborhood.
- range  $\sigma_r$ : “minimum” amplitude of an edge

# Influence of Pixels

Only pixels close in space and in range are considered.



# Varying the Range Parameter



input

$\sigma_r = 0.1$



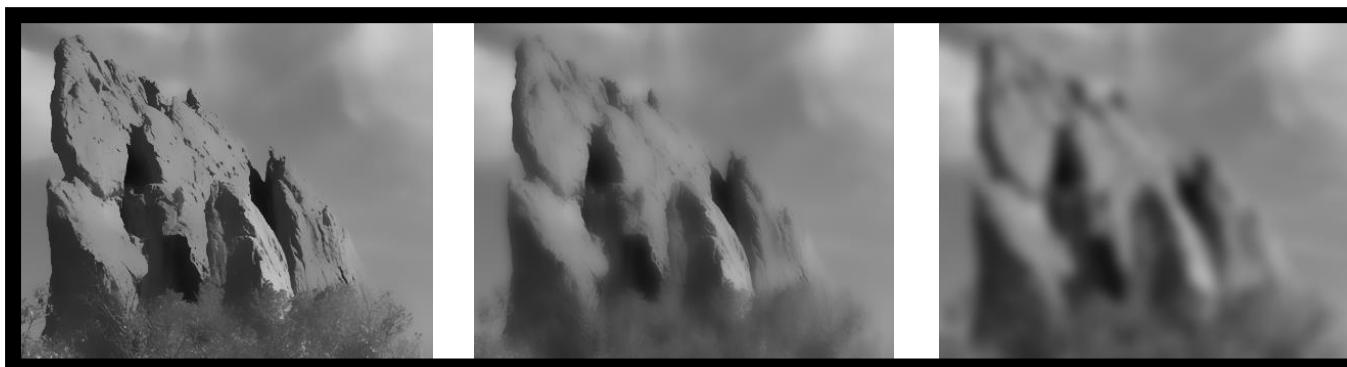
$\sigma_r = 0.25$



$\sigma_r = \infty$   
(Gaussian blur)

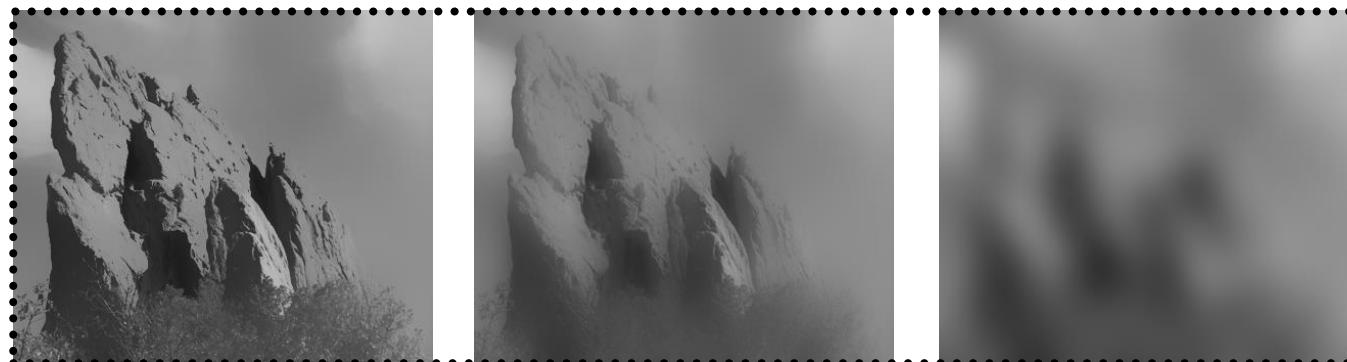


$\sigma_s = 2$



$\sigma_s = 6$

$\sigma_s = 18$

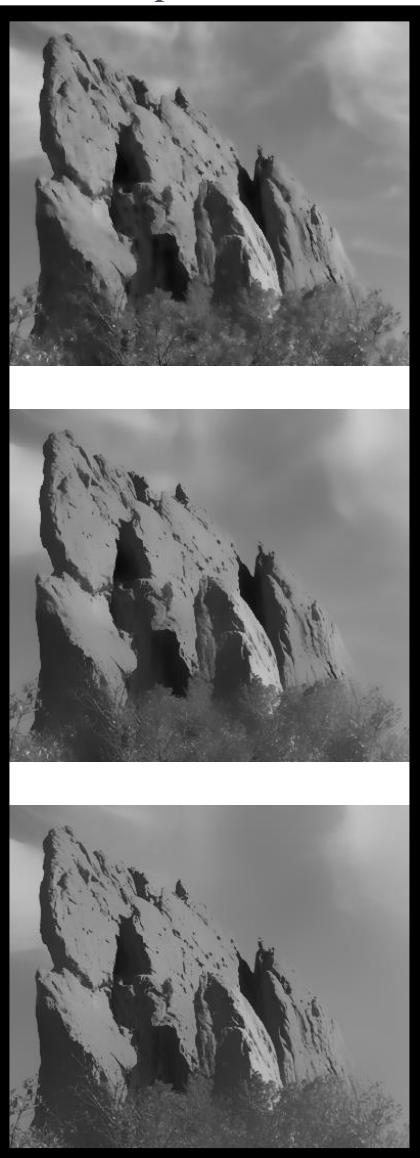


# Varying the Space Parameter



input

$\sigma_s = 2$



$\sigma_s = 6$



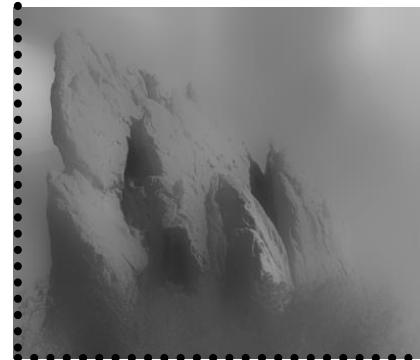
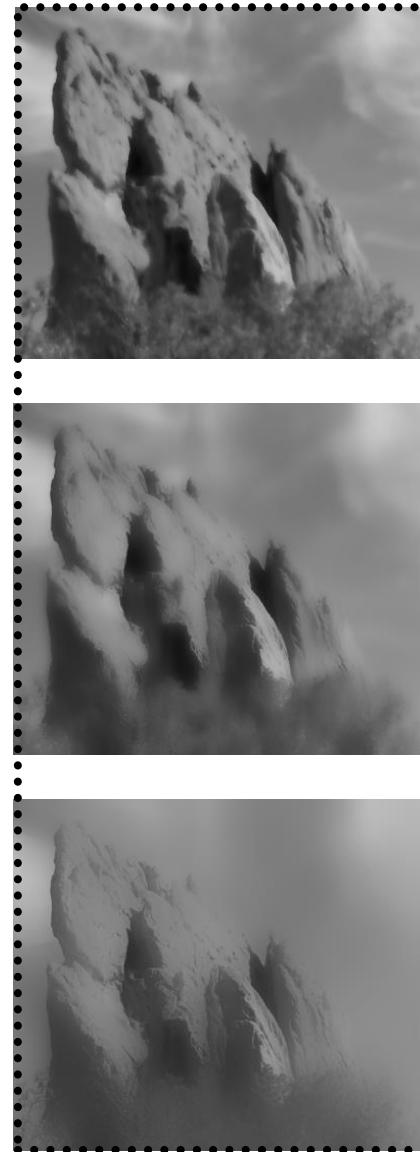
$\sigma_s = 18$



$\sigma_r = 0.1$

$\sigma_r = 0.25$

$\sigma_r = \infty$   
(Gaussian blur)



# How to Set the Parameters

Depends on the application. For instance:

- space parameter: proportional to image size
  - e.g., 2% of image diagonal
- range parameter: proportional to edge amplitude
  - e.g., mean or median of image gradients
- independent of resolution and exposure

# A Few More Advanced Remarks

# Bilateral Filter Crosses Thin Lines

- Desirable for smoothing: more pixels = more robust
- Different from diffusion that stops at thin lines
- Bilateral filter averages across features thinner than  $\sim 2\sigma_s$



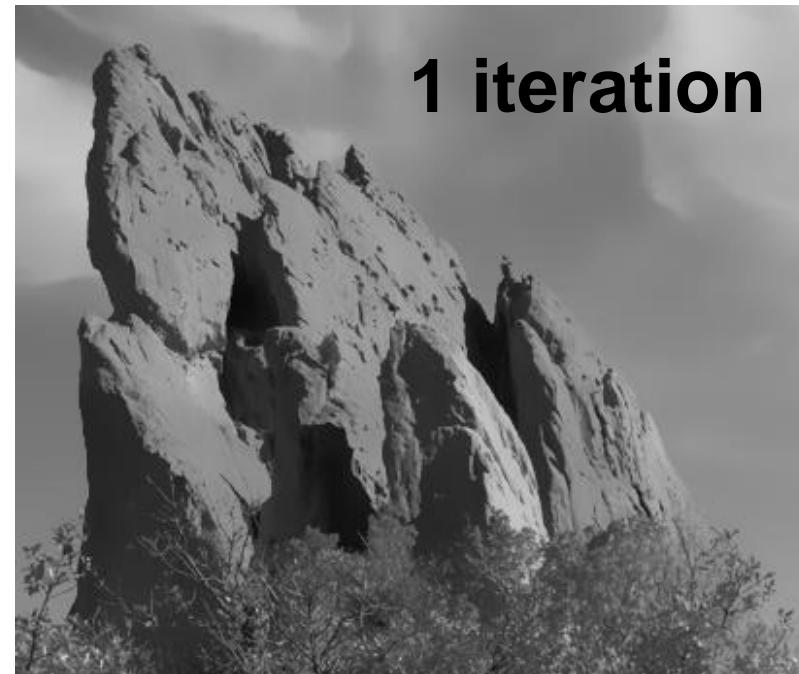
# Iterating the Bilateral Filter

$$I_{(n+1)} = BF [I_{(n)}]$$

- Generate more piecewise-flat images
- Often not needed in computational photo.



**input**



**1 iteration**



**2 iterations**



**4 iterations**

# Bilateral Filtering Color Images

For gray-level images

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(| | p - q | |) G_{\sigma_r}(| I_p - I_q |) I_q$$

intensity difference  
scalar

For color images

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(| | p - q | |) G_{\sigma_r}(| | C_p - C_q | |) C_q$$

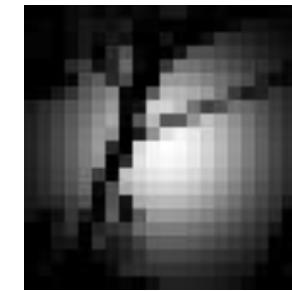
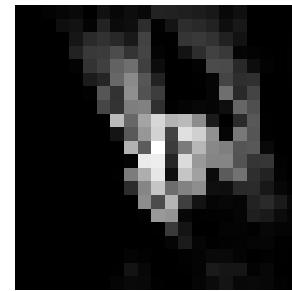
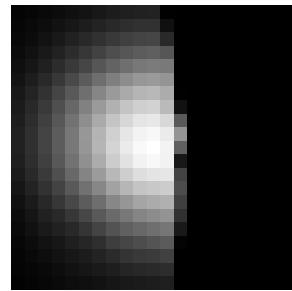
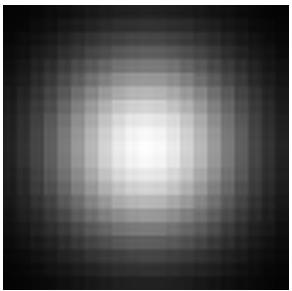
color difference  
3D vector  
(RGB, Lab)



**The bilateral filter is  
extremely easy to adapt to your need.**

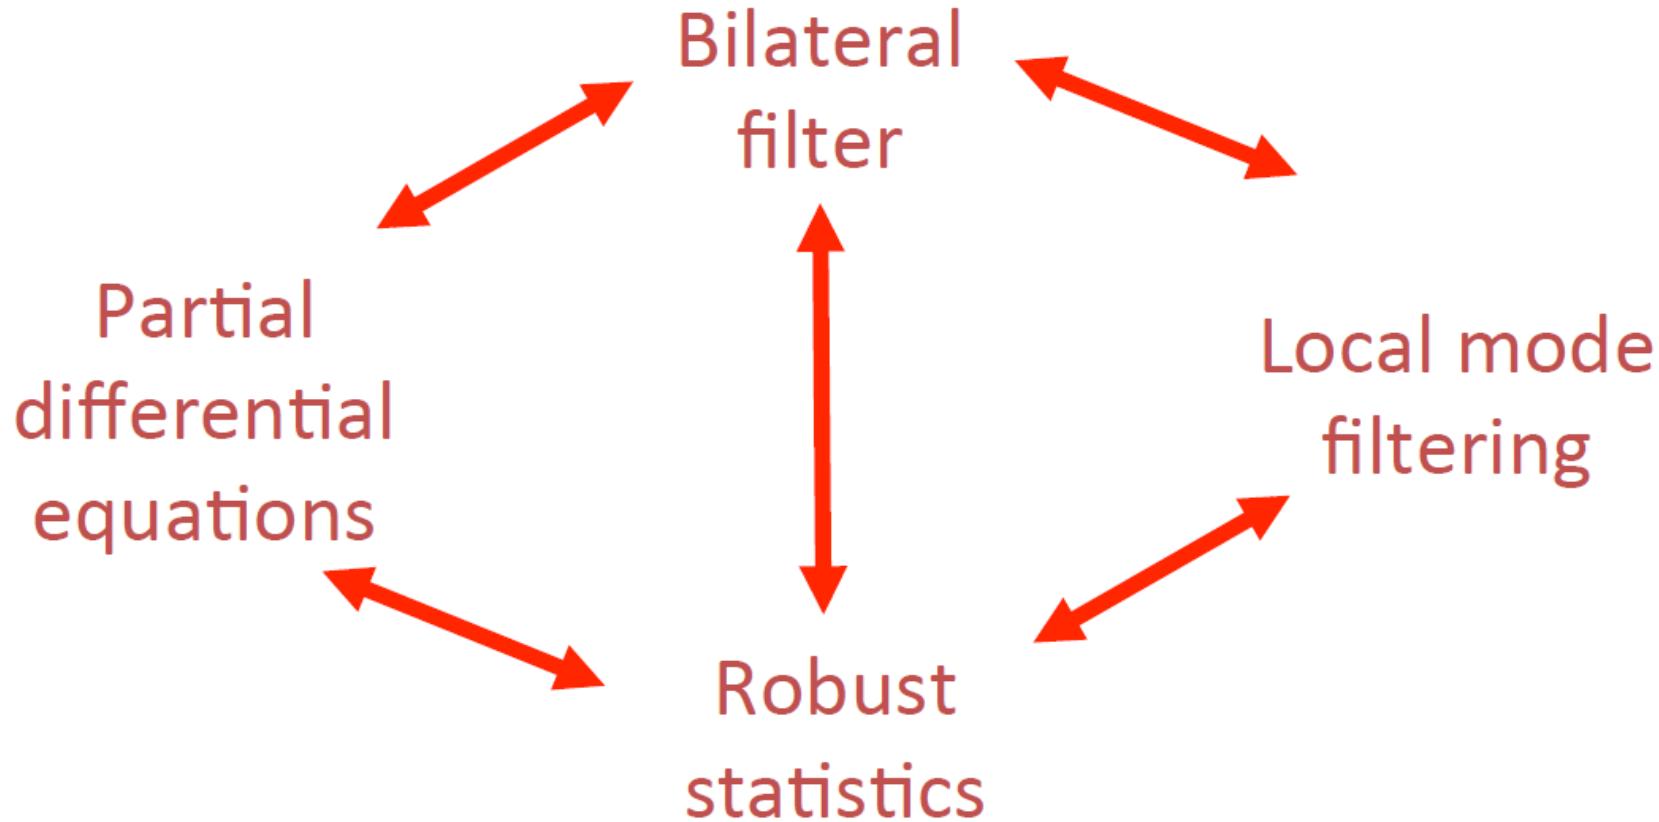
# Hard to Compute

- Nonlinear 
$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(|p - q|) G_{\sigma_r}(|I_p - I_q|) I_q$$
- Complex, spatially varying kernels
  - Cannot be precomputed, no FFT...



- Brute-force implementation is slow > 10min

# Goal: Understand how does bilateral filter relates with other methods



**Additional Reading:** Generalised Nonlocal Image Smoothing,  
L. Pizarro, P. Mrazek, S. Didas, S. Grewenig and J. Weickert, IJCV, 2010

# **Data-driven image processing: “Image manipulation by example”**

**(main idea: pixel patterns in another part of the image are hints for how to improve image in the current region)**

# New Idea: NL-Means Filter (Buades 2005)

- Same goals: ‘Smooth within Similar Regions’
- **KEY INSIGHT:** Generalize, extend ‘Similarity’
  - **Bilateral:**
    - Averages local neighbors with **similar intensities**;
  - **NL-Means:**
    - Averages nonlocal neighbors with **similar neighborhoods!**

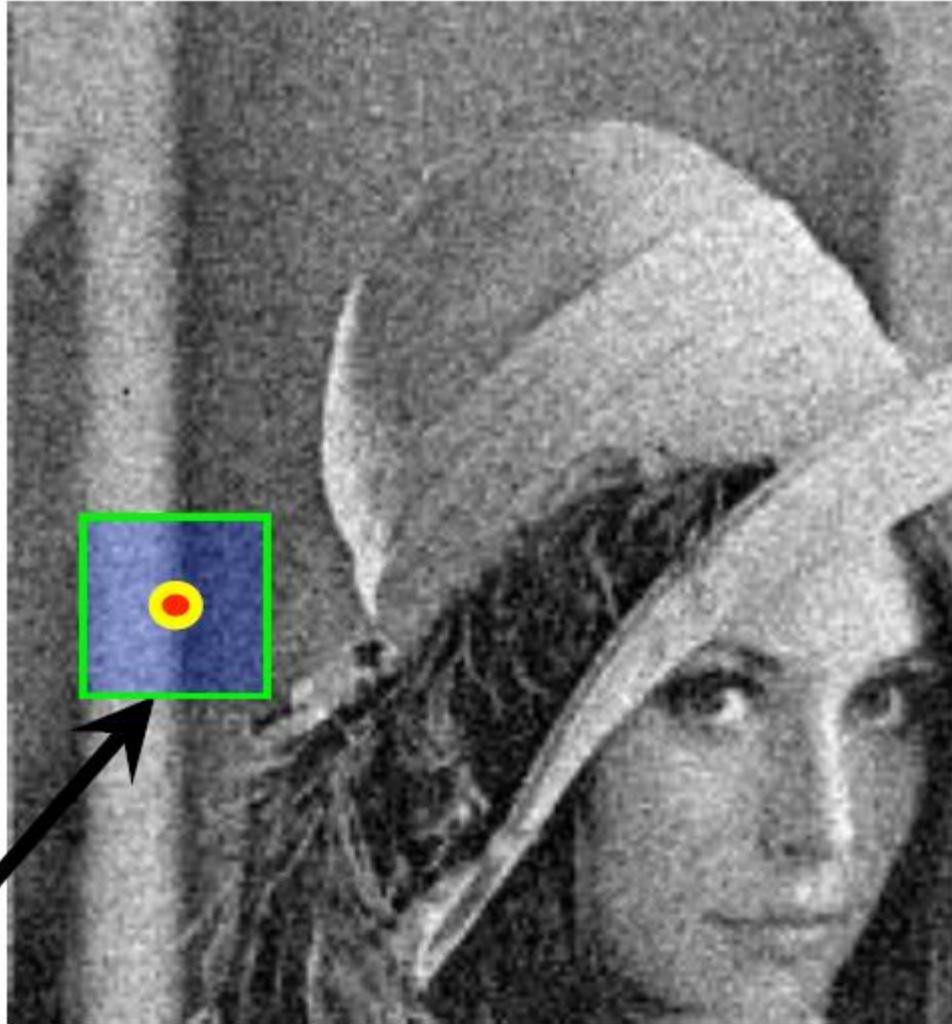
## NL-Means Method: Buades (2005)

- For each and every pixel **p**:



## NL-Means Method: Buades (2005)

- For each and every pixel **p**:
  - Define a small, simple fixed size neighborhood;



# NL-Means Method: Buades (2005)

$$\mathbf{v}_p = \begin{bmatrix} 0.74 \\ 0.32 \\ 0.41 \\ 0.55 \\ \dots \\ \dots \\ \dots \end{bmatrix}$$

- For each and every pixel **p**:



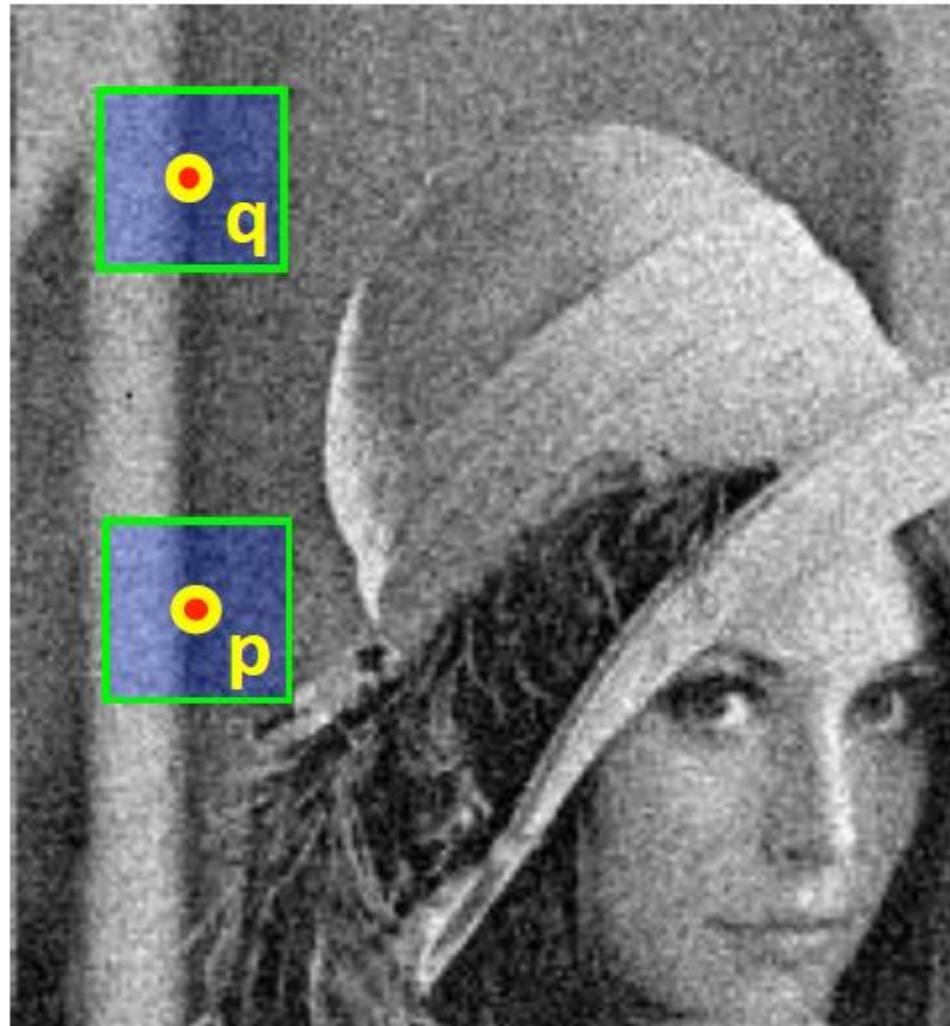
- Define a small, simple fixed size neighborhood;
- Define vector  $\mathbf{v}_p$ : a list of neighboring pixel values.

## NL-Means Method: Buades (2005)

‘Similar’ pixels **p, q**  
→ **SMALL**

vector distance;

$$\| \mathbf{v}_p - \mathbf{v}_q \| {}^2$$



# Denoising using non-local means

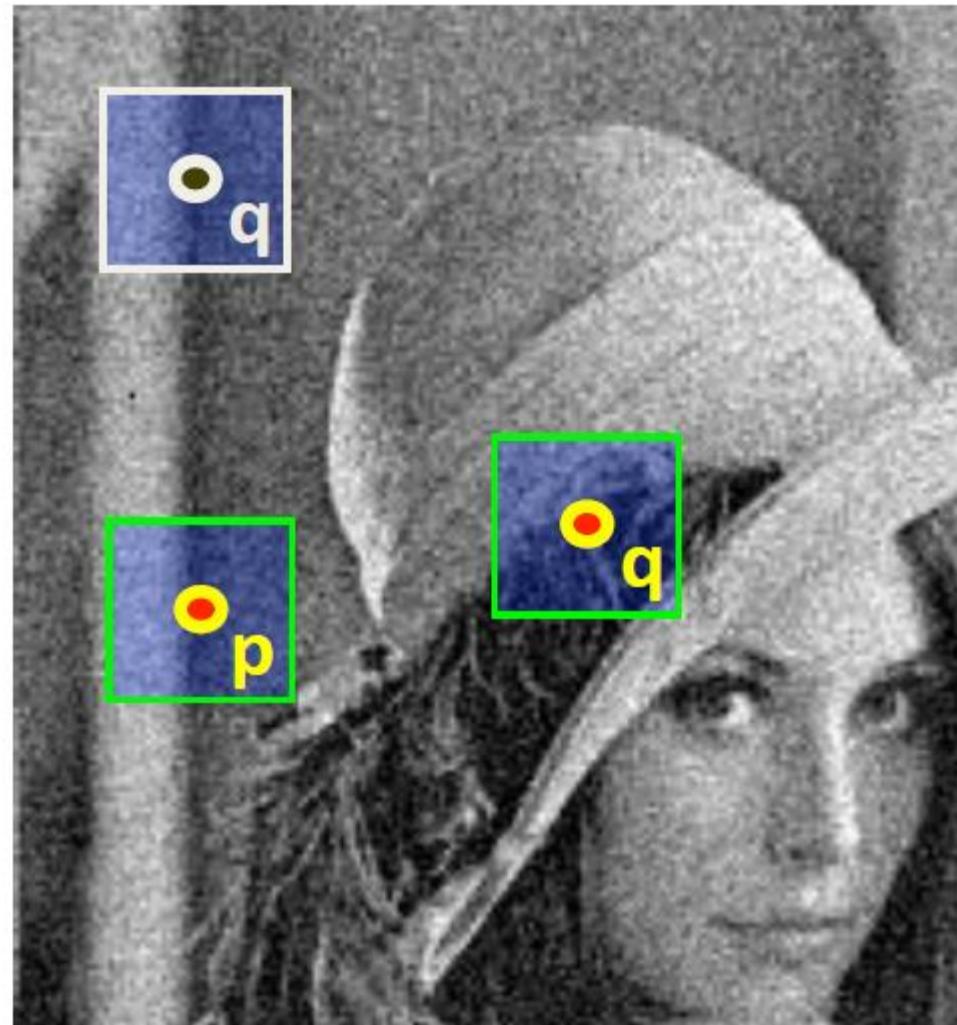
NL-Means Method:  
Buades (2005)

'Dissimilar' pixels **p, q**

→ **LARGE**

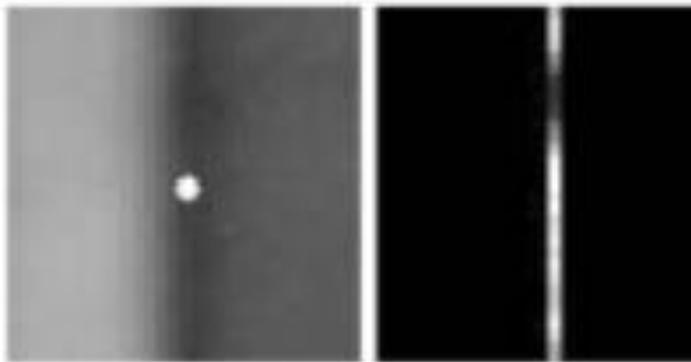
vector distance;

$$\| \mathbf{v}_p - \mathbf{v}_q \| {}^2$$

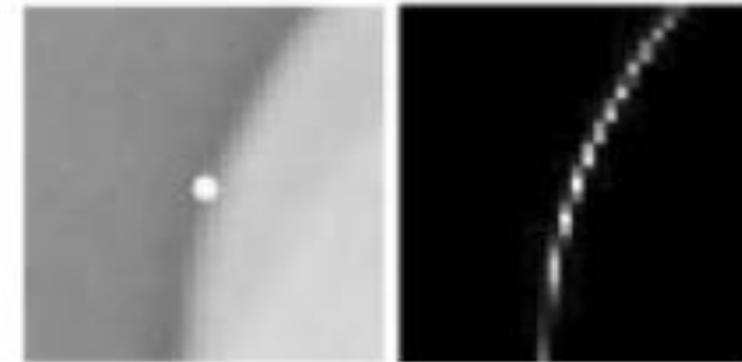


# Denoising using non-local means

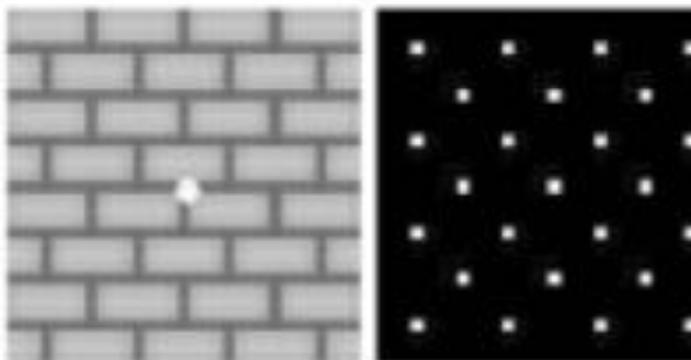
- In each image pair below:
  - Image at left shows the pixel  $p$  to denoise.
  - Image at right shows weights of pixels in 21x21-pixel kernel support window surrounding  $p$ .



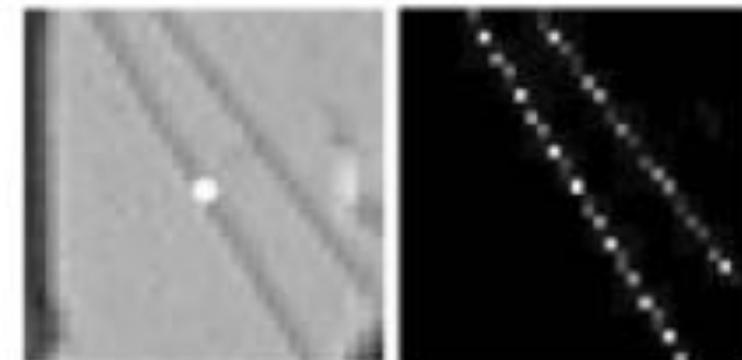
(A)



(B)



(C)



(D)

# NL-Means Method: Buades (2005)

**p, q** neighbors define  
a vector distance;

$$\| \mathbf{v}_p - \mathbf{v}_q \|^2$$

**Filter with this:**

No spatial term!

$$NLMF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\| p - q \|) G_{\sigma_r}(\| \vec{V}_p - \vec{V}_q \|^2) I_q$$



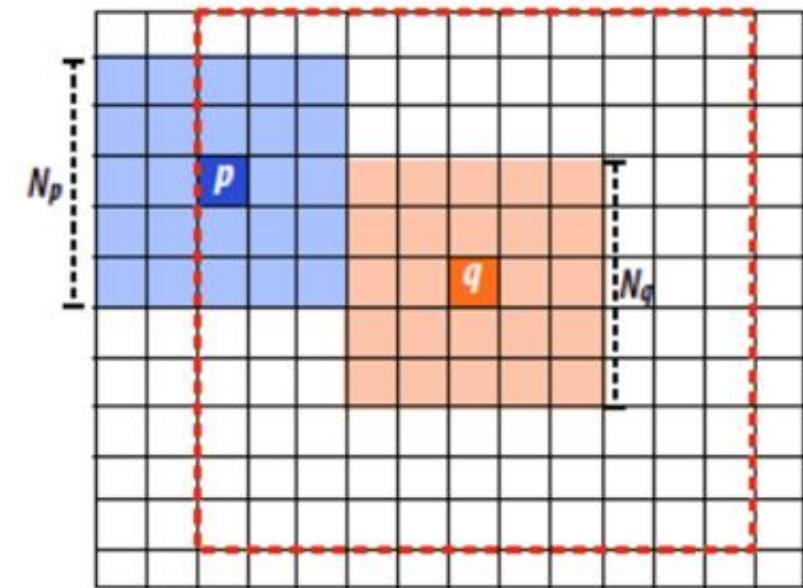
# Denoising using non-local means

- Main idea: replace pixel with average value of nearby pixels that have a similar surrounding region.
  - Assumption: images have repeating structure

$$NL[I](p) = \sum_{q \in S(p)} w(p, q)I(q)$$

All points in search region about p

$$w(p, q) = \frac{1}{C_p} e^{-\frac{\|N_p - N_q\|^2}{h^2}}$$



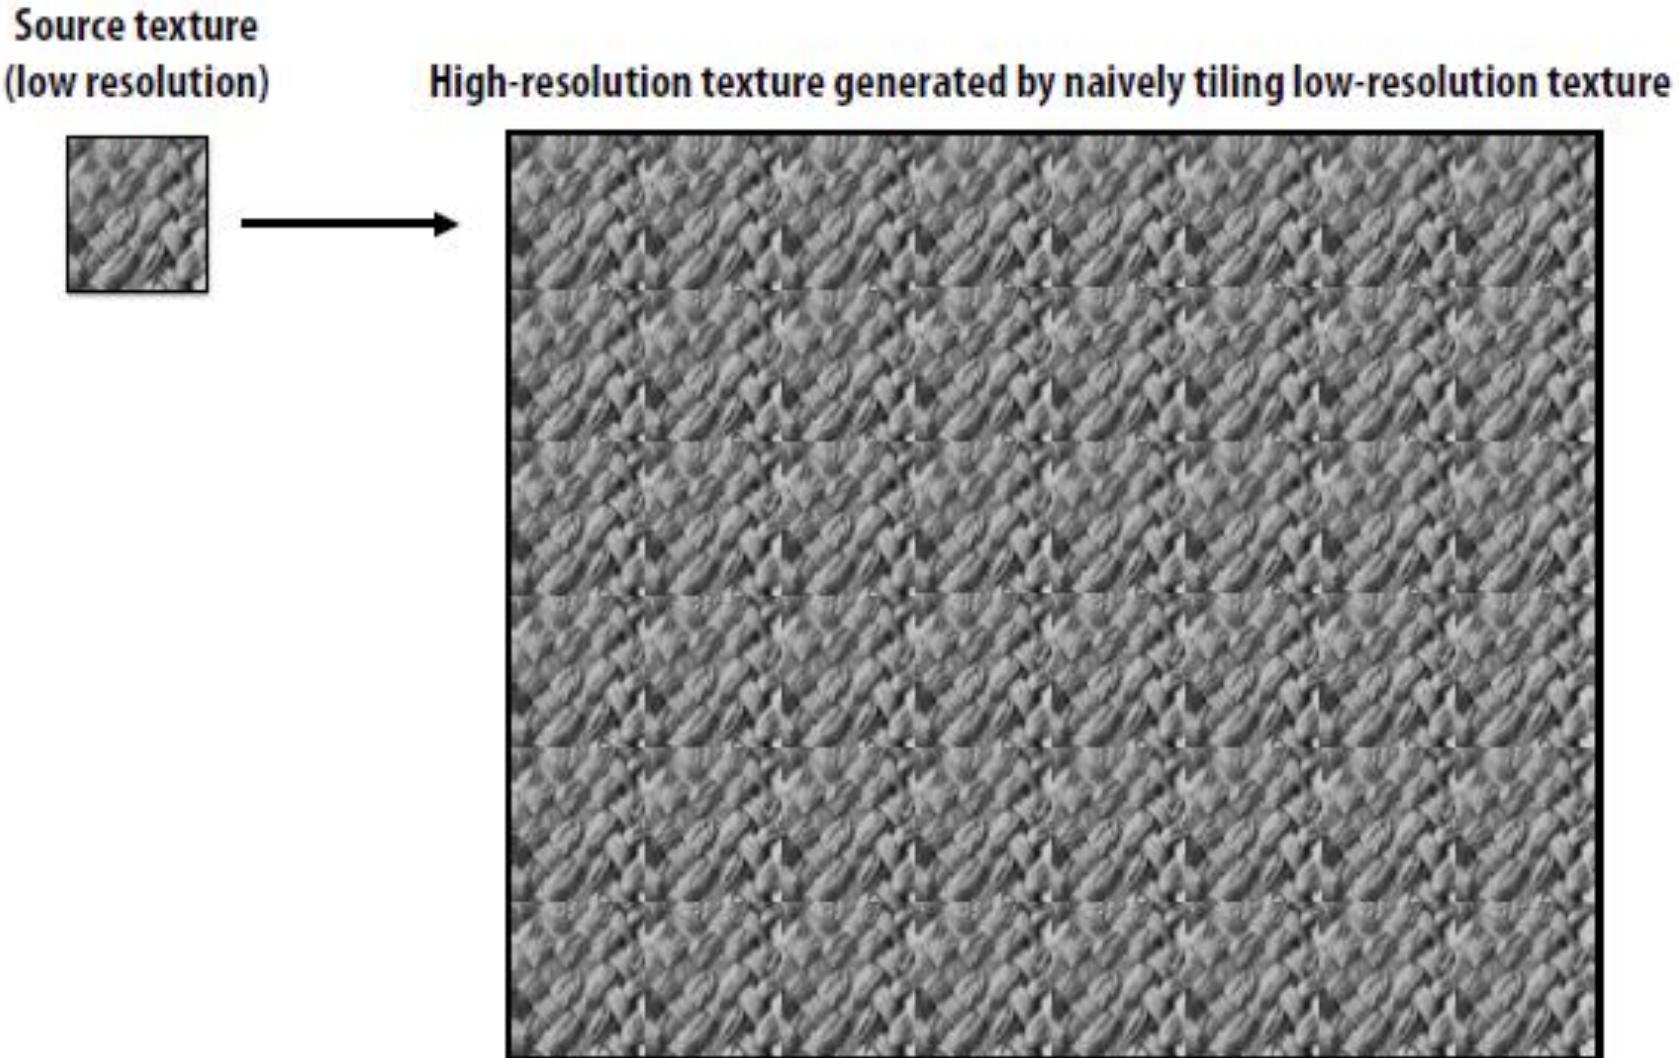
- $N_p$  and  $N_q$  are vectors of pixel values in square window around pixels  $p$  and  $q$  (highlighted regions in figure)
- $L_2$  difference between  $N_p$  and  $N_q$  = “similarity” of surrounding regions
- $C_p$  is just a normalization constant to ensure weights sum to one for pixel  $p$ .
- $S$  is the search region around  $p$  (given by dotted red line in figure)



- Input,
- Gaussian,
- Anisotropic Diffusion, Biliteral,
- NLM

# Texture synthesis

- Input: low-resolution texture image
- Desired output: high-resolution texture that appears “like” the input

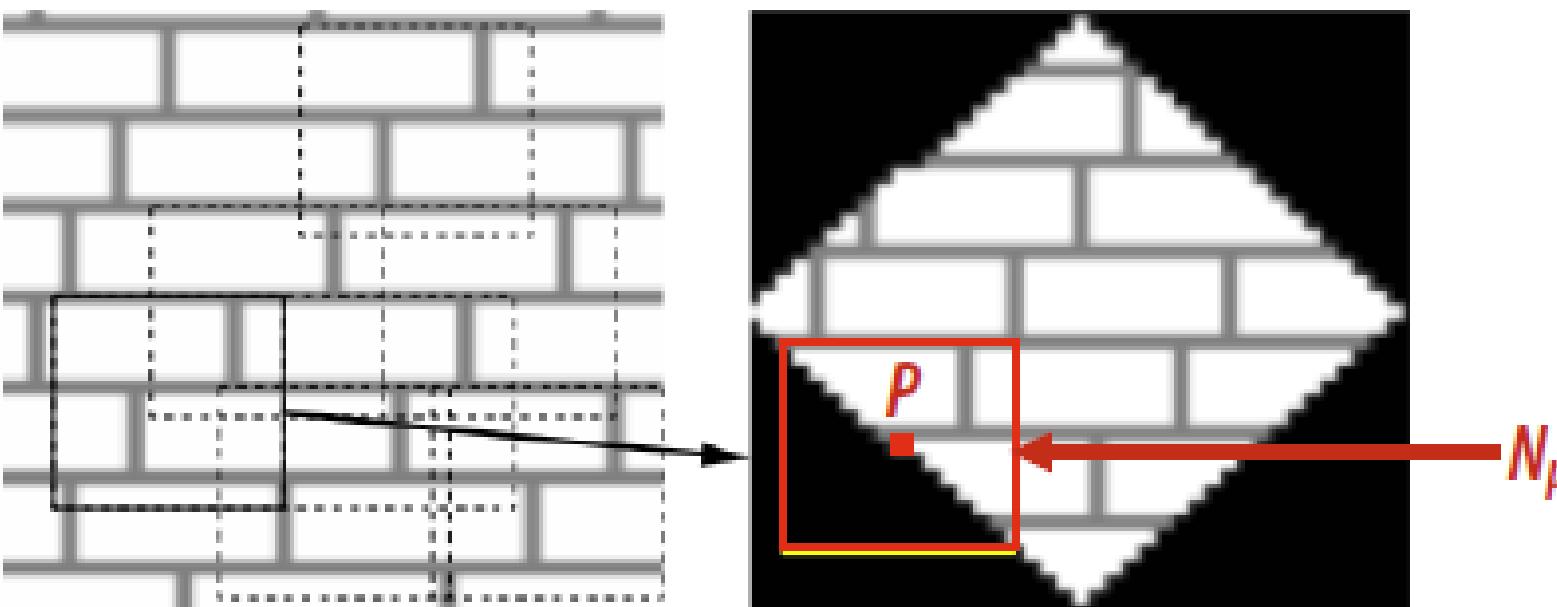


# Algorithm: non-parametric texture synthesis

- Main idea: given the  $N \times N$  neighborhood  $N_p$  around unknown pixel  $p$ , want a probability distribution function for possible values of  $p$ , given  $N_p$ :  $P(p=X | N_p)$

For each pixel  $p$  to synthesize:

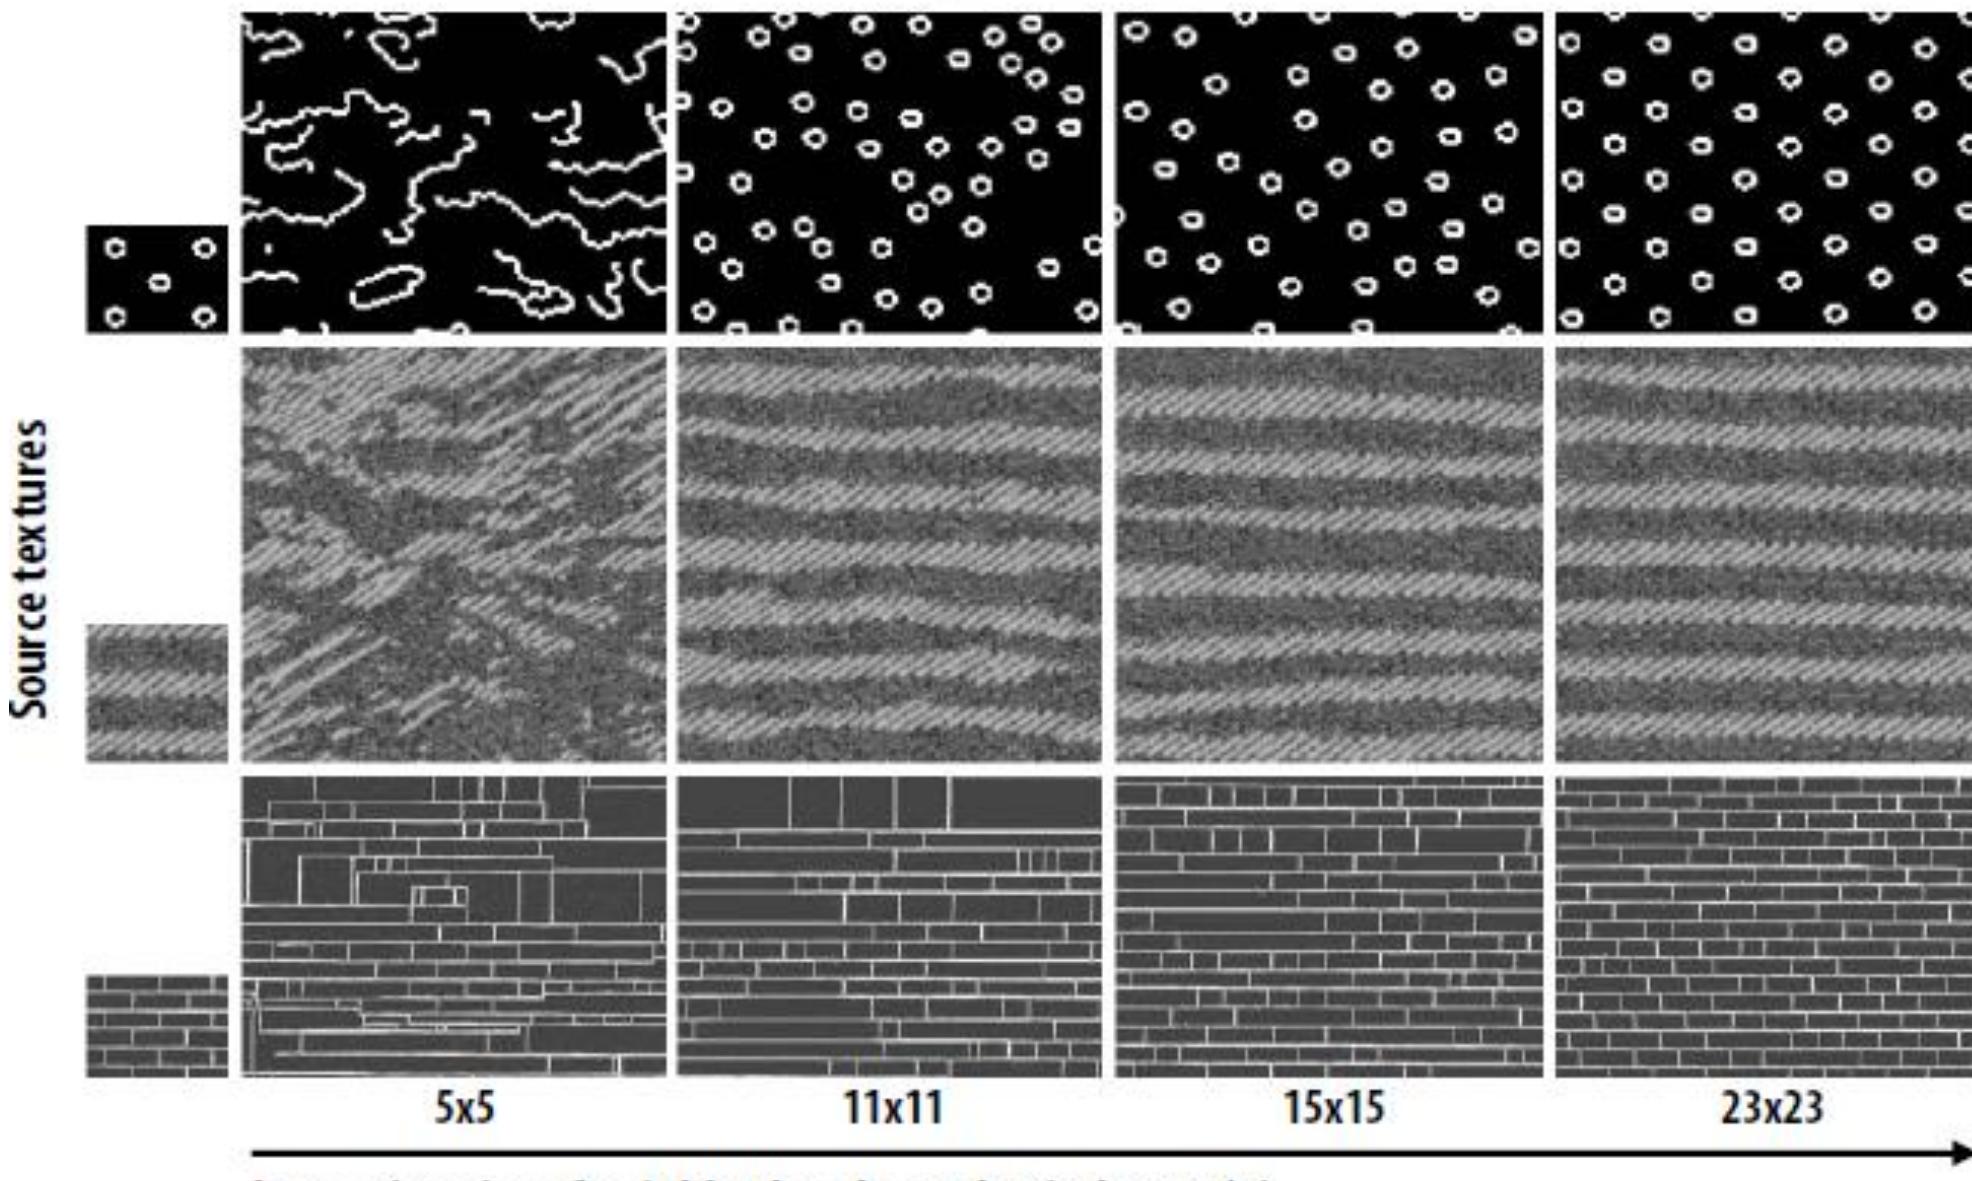
1. Find other  $N \times N$  patches ( $N_q$ ) in the image that are most similar to  $N_p$
2. Center pixels of the closest patches are candidates for  $p$
3. Randomly sample from candidates weighted by distance  $d(N_p, N_q)$



[Efros and Leung 99]

# Non-parametric texture synthesis [Efros and Leung 99]

Synthesized Textures



# More texture synthesis examples [Efros and Leung 99]

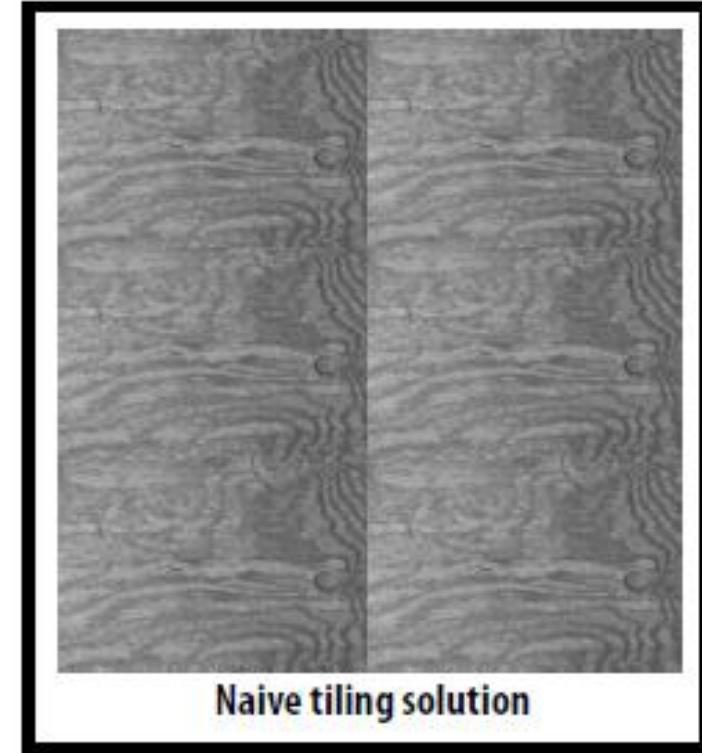


Synthesized Textures



ut it becomes harder to last round itself, at "this daily living rooms," as House De described it last fall. He fall. At the left a ringing question more years of Monica Lewininda Tripp?" That now see Political comedian Al Frat ext phase of the story will

ut it becomes harder to last round itself, at "this daily living rooms," as House De described it last fall. He fall. At the left a ringing question more years of Monica Lewininda Tripp?" That now see Political comedian Al Frat ext phase of the story will



# Image completion example



Original Image



Masked Region



Completion Result

**Goal: fill in masked region with “plausible” pixel values.**

See PatchMatch algorithm [Barnes 2009] for a fast randomized algorithm for finding similar patches

Image credit: [Barnes et al. 2009]

# Image processing summary

- **Image processing via convolution**
  - Different operations specified by changing weights of convolution kernel
  - Separable filters lend themselves to efficiency implementation as multiple 1D filters
- **Data-driven image processing techniques**
  - Key idea: use examples from other places in the image as priors to determine how to manipulate image
- To learn more: consider CMU 15-463/663: “Computational Photography”

Any questions?



# References - courses

- “A Gentle Introduction to Bilateral Filtering and its Applications” given by **Sylvain Paris**, Pierre Kornprobst, Jack Tumblin, and Frédo Durand ([http://people.csail.mit.edu/sparis/bf\\_course/](http://people.csail.mit.edu/sparis/bf_course/))
- Bilateral Filtering, and Non-local Means Denoising, **Erkut Erdem**

# References

- Siggraph 15: An L1 Image Transform for Edge-Preserving Smoothing and Scene-Level Intrinsic Decomposition. [code]
- Siggraph 14: Bilateral texture filtering
- Siggraph 14: Fast Local Laplacian Filters: Theory and Applications. [code]